



UNIVERSITÀ
DEGLI STUDI
DI PADOVA



INFORMATION ENGINEERING DEPARTMENT
COMPUTER ENGINEERING - AI & ROBOTICS

Intelligent Robotics 2023/2024
Francesco Crisci 2076739

Introduction

Lecture 3

One possible exam question is the DOF of a synchro drive robot and why it is that. another question might be about robot exploration, mapping and localization. (Do not trust google to study or the exam).

Explain how tactile skin for a robot can be constructed, elements and how they are assembled. (This professor is basically an asshole about questions, so be as specific as possible and explain everything that is in the book and slides).

Some historical shit no-one cares about.

Automation and Autonomy

There is no intelligence without a body. we'll see it later on.

Autonomy means dealing with unknown problems. What is an intelligence robot? In Ai, a robot is intelligent if it can achieve is goal. AI typically is divided in 7 major area: knowledge representation, learning, learning, problem solving, natural language processing, inference, search and vision.

Autonomy: The quality or state of being self-governing, self-directing freedom and especially moral independence. We mean self-governing as the James Watt self-governor. It means taking action depending on the information from the outer world. Bounded rationality means that the capability of AI to have rational thinking is really limited. Strong AI can be applied to multiple domains, weak AI can be used to a single domain. Bounded rationality derives from the limits of cognitive agents. Robot cannot get drunk like us, so we win. Autonomy is the capability of generating action according to what happens in the environment. To take an action, or in general to be able to complete the goal, it is different from cognition. Autonomy is the capability of reasoning about a modelled event.

It is about a physically-situated agents who not only perform actions but can also adapt to the open world when the environment and tasks are not known a priori by generating new plans, monitoring and changing plans, and learning within the constraints of their bounded rationality.

Automation is about physically-situated tools performing highly repetitive, pre-planned actions for well-modeled tasks under the closed world assumption. It works in the close world assumption, i.e., we know every detail about that world.

Industrial robots

We have a mature technology, with strong basis of control theory, faster joint movement. They are high repeatable. Thou, they need fixturing. A fixture is a support, or reference, that you have to insert in the working space of the robot so that you ensure that an object can be manipulated.

a fixture keeps the item in the correct position at the right time. They are now adding sensors to reduce the fixtures, in order to add some flexibility in the case the fixture was not in the right place.

world assumption

The close world assumption implies that everything in the robot environment is known or controlled. Everything is known a priori, this means you can use formal logic to reason about the world, since we can substitute objects with symbols. Everything can be completely modeled and we can create exiting procedure if there are unexpected situations. We can also minimize or eliminate sensing. You can Engineer the world(He really likes this word so i'll use it in the exam), i.e, you change the world in order to simplify your problem/problem.

Lecture 4

We want to describe the difference between automation and autonomy; understand how this difference has shaped different cultures within robotics; define...

Automation is usually related to industrial robots working engineered environments. Always repeating the same actions. There is some intelligence as sensors, segmentation, localizing algorithm and so on. we know everything about the world. For autonomy we usually can think about rovers, which they explore a non closed world, where they do not know if the soil is solid or not, if there are martian, water and so on. We just know so many little things about this environment and the robot has to deal with the unforeseen. An autonomous robot works in an open world. We cannot assume what is happening or what the robot encounters. These two assumptions define close world and open world.

Close world

You know the environment, so you can create a sequence of actions in which we can assume the success of all actions since we know what is gonna happen.. This means you can create a plan based on formal logic, based on sensors to use them as control signals or use them directly to trigger a machine. The focus is to create a stable control loop to respond to all expected situations. This includes also failures, not only the correct situations. This helps in minimizing the sense of the robot, if not eliminating it. This assumption has many problems, like the toy problem called the monkey and the banana (frame problem, because you have your world and you know an area of your world, and you are trying to model everything inside that frame, but you do not know what is happening outside of this frame). If we only manipulate symbols (mapping objects of the real world into symbols) we are missing the fact that there might be some configuration we aren't able to model or we didn't model. There is additional information that we cannot put in the modelling of the world. It comes from the fact that the real world is never a closed world. You never have the whole knowledge on it.

Open world

we do not make assumption that everything can be modelled, or that we can predict all the actions. We assume from the beginning that we have a partial knowledge from the environment, and we assume that this knowledge is only partially correct. There might be some error, things we cannot

perceive. We try to design our robot in order to be compliant to this world. We put enough sensing in our robot so that it can understand the world at the best. We need to have some system to dynamically change the actions of the robot. We need to cope with actions and the world. You cannot model everything in the world. This is associated with being autonomous. the difference between autonomy and automation affects the hardware design (think of vacuum cleaner and an industrial arm), the programming style and the kind of failures.

Why it matters?

It affects hardware design: autonomy requires rich sensing in order to monitor the key elements of the dynamic world, so a robot designed for automation or for teleoperation is not necessarily able to be used autonomously by adding software. Think of it as an ecology, what ecological niche is the robot going to fill.

It affects programming style: if you have a closed world, delegating for a small set of repetitious tasks, then we focus on formal, stable control loops. On the other hand, if you focus on open world, delegating for a variety of tasks while operating in dynamic environment, we focus on artificial intelligence. The first case is automation, the second is autonomy. For the first case we can think of military planes, such as drones and stuff. for the second case, you can think of robot puppies, which can do stupid things but are way cuter eheheh.

It affects how system breaks: which kind of failure you can expect from the system. a machine cannot substitute a person, because there are things that are not correctly modelled, and so if there is a person thinking about a situation he can recover from the unknown but a machine cannot. Thinking of devices that are controlled by humans and then you want a machine doing things by itself. This happened in the military, where they moved to almost autonomous flights or vehicles, controlled by remote operators.

Can you solve a problem with automation or autonomy?

Automation

- Execution: perhaps plan once, then repeat that plan forever
- deterministic: can model the system deterministically
- Closed world: the model contains everything
- Signals: control or decision making is at the signal level

Autonomy

- Generate: constantly generating new plans
- Non-deterministic: system is too complex to model deterministically
- Open world: models will only be partial
- Symbols: control or decision-making is with symbols or labels.

The path to autonomy is never guaranteed, but sometimes it is necessary. There are some applications which mix the two of them, like which part of the problem can be solved with automation and which part with autonomy. The robot vacuum cleaner has the following slider: the plan slider

goes more towards the execution since the path is always the same and does not change through time. The action plans it goes more through the non-deterministic since it is always checking its own position. It continuously try to localize itself. The models goes towards more the close world assumption than through the open world. The knowledge representation goes towards the signals, but the most advanced ones use some symbols too.

Locomotion

There is no brain/intelligence without a body. You cannot have two different teams working separately on brain and body. It would not work.

Embodied intelligence: the intelligence is not only in the brain but is also in the body. You need some intelligence in the mechanical design.

When designing a robot you should put intelligence also in the body.

Locomotion from Class notes

We have 5 different type of locomotion:

- wheeled
- legged
- Snake
- Free-floating
- Swimming

Mobile robot kinematics is about models-maneuverability-motion control. We will focus on wheeled locomotion. Locomotion can be achieved with different mechanism and it involves different aspects: Actuators, properly designed, Forces control and the reactions to the forces, managed by control theory systems, dynamics, passive dynamics and animal locomotion to understand and get inspired for our robots. We can have hydraulic or pneumatic muscles for some types of robots to achieve movement. We need to consider the following aspects: Friction and direction, design and control, Actuation and sensing, animal and machine. In order to move in the environment we need Stability (number of contact points, center of gravity, static vs dynamic stabilization, inclination to terrain), Contact (contact point or area, angle of contact Friction) and Environment (structure and medium. You cannot project a swimming robot and put it in the desert.). A robot to be stable it need at least 3 contact points. The stability can be challenged from the environment, such as inclination of the terrain and so on. The number of contact points is not the only issue, we need to consider also the angle of contact of the contact points. There is an important table on the book, hopefully it is on the slides.(It is in. lesson 4).

The concept found in nature are difficulty to recreate artificially.

Legged locomotion

It is nature inspired. The movement of walking biped is close to rolling. The number of legs determines stability of locomotion. Bipedal locomotion is really hard to recreate, consider that a

human learns it in 2 years or so. Walking of a biped is not too far from real rolling. Rolling of a pol....

Legged systems can overcome many obstacles that are not reachable by wheeled systems. But it is quite hard to achieve this since many DOFs must be controlled in a coordinated way and the robot must see detailed elements of the terrain. The wheeled robots cannot overcome obstacles that have an height greater or equal than the radius of the wheel. Wheeled robots are worse than legged robots for stairs. Legged robots are more stable than wheeled ones. Wheeled have problems with slopes and discontinuities, while legged robots an overcome them. In some cases, like short discontinuities in the terrain, the legged robots can have some problems, like getting stuck or something. when designing, should we consider walking or rolling? You have to consider the following: How many actuators you need, the structural complexity, the control expense, the energy efficiency, thee terrain (flat ground, soft ground, climbing,...) and movement of the involved masses (walking/running includes up and down movement of COG(center of gravity); some extra losses). Crawling/sliding is requires more power compare to running, which is more consuming compared to walking, which is more consuming than railway wheels. But the top speed can be really different... Remember that when we are walking we are moving the center of mass, so we should consider exploiting this property. But how can we generate motions for walking? We have different ways:

- Manual motion design: classical slider based motion editors. You create different positions (like frames of an animation). We design different positions configurations and then we can achieve an almost fluent walk for the robot. Unfortunately, it takes a lot of time and it is trial and error procedure.
- Control based: ZMP(zero moments point: technique to make humanoid robots walking) and inverse pendulum based control.
- Learning algorithms: motion is specified in a parametric way(Central pattern generation); the parameters are determined by search algorithms(Genetic algorithms,...).
- Motions are obtained from human data: learning by watching/imitation; motion re-targeting. Usually the kinematic chain of a robot is different from the one of a human, that is why we use motion re-targeting, i.e., re-projecting the human kinematic chain(human motion) on the robot kinematic chain. It is used in game/movie industries; We use an actor and then a motion capture system to recreate some movements.

This is just about finding the position of the joints, we are not considering the dynamics and the structure of our robot, which is really important for the energy efficiency of our robots.

A passive walker is a robot with limbs that has no motors. You exploit gravity and structure of the legs that have springs in the structure which release energy during the walking. The energy comes from either the kinetic energy (like pushing the robot) converted in elastic energy either potential energy, like going downhill and it is transforming the potential energy in kinetic energy to move. Of course, it cannot move uphill. Asimo walks with the knees bent, which makes the leg extension process more easy since you do not move the center o gravity but you expend a lot of energy, so it is energy inefficient. We want to combine passive walking with dynamic stability. The study of passive walker...(Go to slides). We define a metric for the efficiency of a walker as

$$COT = c_{mt} = |mech.energy|/(weight \times distancetraveled)[JKg^{-1}km^{-1}]$$

a clever design o the legs can be really useful. Static walking can be represented as an inverse pendulum.

The fewer legs the more complicated becomes locomotion: stability, at least three legs are required for static stability: babies have to learn for quite a while until they are able to stand or even walk on their two legs. For static walking at least 6 legs are required. during walking some legs are lifted.


Wheeled locomotion

We have two types of wheel: Standard wheel with 2 DOF (It can go straight or if you have one wheel you can turn around in the same spot. It can turn around the point of contact. Think of an acrobat on the unicycle) or castor wheel with 3 DOF (It has two axes: one to spin the wheel and one to the point the wheel is attached on the chassis. We have one DOF for the spinning, one for rotating around the point of contact and one for the rotation about the contact point of the chassis). The standard wheel is fixed to the frame of the vehicle. Chassis of the ROBOT is the frame (or the box) composing the body of the robot and it is the structure to which we attach the wheel. We want the axis of the wheel to be fixed on the chassis of the robot. We have also other types of wheels: the Swedish wheel with 3 DOF (aka Mecanum wheel or Omnidirectional wheel. It can move also sideways. It can have rotation along the wheel axis, around the point of contact and it can move sideways. That is why it has 3 DOF. The first two are the same for the standard wheel) and spherical wheel with 3 DOF (You have a sphere encapsulated in a support attached to the chassis of the robot, like the one you have in old mouse for computers).

Locomotion again

It is the act of moving from place to place: we have seen synchro drive with 2 DOFs and omnidirectional robot with 3 DOFs. Synchro drive cannot generate torque. 4 wheel robot needs a suspension system. We can also have 6 wheels arrangements. We can have platforms to help with the stability of the robot (platforms instead of wheels). Remember that a wheel, to overcome an obstacle, needs a radius that is bigger than the height of the obstacle itself. If you think of the Shrimp 3rd passively climbs the obstacles that are higher than the wheel radius. It is a 6 wheel configuration robot. This structure is compliant (it can adapt) to the morphology of the terrain. Locomotion relies on the physical interaction between the vehicle and its environment. Locomotion is concerned with the interaction forces, along with the mechanism and actuators that generate them.

Motion and coordinate frames





We can reduce the robot to a point and to a vector. Robot is a vector in the 2D Cartesian plane in the 2D case. We have 2 coordinate frames:

- world frame $W(x, y, \theta)$. We have 3 DOFs, aka 3 variables we can change to manipulate the robot
- Robot's reference frame R, useful to move around and give them perception since they should be intelligent

In the 3D case we have the same things, W and R, but in 3D we have 6 DOFs of motion. In 3D perception and actuation are in the robot's reference frame. The goal is to find the correct transformation between W and R. If we have a robot with 1 DOF it means that it can move only in one direction; the robot can translate (or rotate) along 1 dimension(axle). We can think of a train as a 1 DOF machine. For 2 DOF we have a robot that can translate(or rotate) along 2 dimension; for example a telescope or the synchro drive robot.

Differential drive

Consider the following 2 figures: By selecting the speed of each wheel we can control the motion



of the robot: if the two wheels have the same speed, the robot will move straight; if the wheels have same speed but opposite direction, the robot will rotate on the spot; if one wheel goes slower than the other, the robot will turn on the side of the slower wheel. You are setting a circular trajectory and an angular speed. Let now consider the following:

- V_L and V_R as the speed of the left and right wheel respectively
- ω_L and ω_R as the angular speeds of the left and right wheels respectively
- W as the distance between the wheels
- $L = R - W/2$ and $R_i = R + W/2$ as the distances of the left and right wheels from the center of the trajectory respectively

We know that that $V_L = \omega_L r_l$ in case there is friction. In case of ice we have that the velocity and the radius are not zero, but the angular velocity of the wheel is zero. If the velocity of the left and right wheel are the same, we have straight motion; if they are the same but opposite in direction, the robot will turn on the spot. If they are different we'll have the following properties:

- Left wheel: in Δt the traveled distance is $V_L \Delta t$; the radius of arc is $R - W/2$
- Right wheel: in Δt the traveled distance is $V_R \Delta t$; the radius of arc is $R + W/2$

We then have: $\Delta\theta = \frac{V_L \Delta t}{R - W/2} = \frac{V_R \Delta t}{R + W/2} \Rightarrow \frac{W}{2}(V_L + V_R) = R(V_R - V_L)$. Then we have the following two equations:

$$R = \frac{W(V_R + V_L)}{2(V_R - V_L)} \quad \Delta\theta = \frac{(V_R - V_L)\Delta t}{W}$$

Where W is a constant and Δt is the time elapsed. The angle is regulated by the difference of the two velocities. For the configurations we have seen the following will happen:

- if $V_R = V_L \Rightarrow \Delta\theta = 0, R = \infty$;
- if $V_R = -V_L$, or, $V_L = -V_R \Rightarrow \Delta\theta = \frac{2V_R \Delta t}{W}, R = 0$

Locomotion from Murphy

10 minutes on why you should take notes and why ghidoni is a chad.

For autonomous robots is important to move around. A mobile robot can move and sense, and must process information to link these two. Possible locomotion systems are speed and/or acceleration of movement; precision positioning; flexibility and robustness in different conditions; efficiency. We need to be able to say if the motion of the robot has been accurate enough, since it might think it moved a meter but actually moved 80cm or 1.20m. You have a distribution of errors leading to an not precise positioning. We also want to cope with locomotion in different conditions, so we want it to be robust and flexible in different situations. We'd also like to program a robot that is efficient enough and with lower power consumption.

Locomotion is the act of moving from place to place. It relies on the physical interaction between the vehicle and its environment. It is concerned with the interaction forces, along with the mechanisms and actuators that generate them. A person has 3 degree of freedom: x,y and theta(the rotation angle to change the direction). To determine the position of a robot you need 3 degrees of freedom, x,y and the orientation angle. a robot can also have 3 degree of freedoms if they can move in the 3-dimensional space and can rotate on all 3 axes. So we have 3 DOF for traslation and 3 for rotation.

Mechanical locomotion

Many approaches i AI robotics assume that the robot is holonomic, which means that the device can be treated as a massless point capable of instantaneous turning in any direction. It allows the designer to ignore the complexity involved in mechanical control and we do not have to worry about the robot slowing down before making a turn or having parallel park. It simplifies the planning and localization. In reality robots are nonholonomic since robots have a mass, and no matter the design, there is always some skitter where the robot crabs to one side when it turns. Due to the fact that robots are nonhomonolic, the reactive or deliberative functions will need to know how fast the robot is moving and how far.

Steering

Ground robots are nonholonomic, so it is important to understand how they are steered. We have two designs that approximate holonomic turning for benign conditions:

- *Synchro-drive*: where the wheels must turn together, approximating turning in place, though there is some slippage of the wheels on the floor. The wheels are linked through either a drive belt or a set of gears. The wheels tend to be narrow to minimize contact with the ground and thus they skitter.
- *Omnidirectional wheels*, also called Mercaum wheels, are wheels with rollers that allow the robot to move perpendicular to the main wheel. By adjusting the relative speeds of the forward and lateral wheels, the robot can go sideways. Unfortunately, these wheels tend to be easily clogged with dirt and gravel in outdoor applications.

There are also two major styles of steering for nonholonomic vehicles, and these are used extensively in real-world applications. These two styles are:

- *Ackerman Steering* is how a car is steered; the wheels in front turn to guide the vehicle. Technically it refers to the mechanism that adjusts the wheels so that the wheel on the inside of the turn is rotated more than the wheel on the outside of the turn in order to compensate for slight difference between the inside and outside wheel diameters. It is truly nonholonomic, considering how hard it is to parallel park.
- *Skid steering*, also known as differential steering, is how a tank or bulldozer is steered. The tracks on each side can be controlled independently. If the robot is to turn left, the left track slows down while the right track speeds up. Skid steering is not restricted to tracks. It can be implemented on wheeled vehicles. It can allow the vehicle to approximate holonomicity but how closely depends on many variables, such as favorable surfaces and power of the platform.

Polymorphic Skid Systems

They are used prevalently in military robots. These robots have tracks that can change shape to adapt to terrain. Consider the following: the larger the area the tracks contact, the more traction and the more likely the track will be able to move and climb. However, the larger the contact area the more power it takes to turn as the robot pivots on the inner track and the inner track must overcome friction to act as a pivot. Thus it would be desirable to have a track system where the amount of contact area could be adjusted for the situation. Polymorphic skid systems allow tracks to have a variable geometry. One approach is to create a tracked system where the continuous track changes shape, from flat to triangular. The advantage is that it is highly adaptive and requires less power. The down side is that, as the track lift up to climb over an obstacle or turn, the lift creates a gap or opening. Object or debris may get into the gap causing the track to separate from the wheels moving the track, eventually disabling the robot. A second disadvantage is that the tracks tend to be flexible and loosely fitted so that they can adapt to different configurations, but the track may come off the rollers, or detrack.

Another approach to polymorphic skid-steering is to add tracked flippers. The robot relies on a pair of continuous flat tracks but can raise and lower tracked flippers to go over obstacles or gain greater traction or maintain balance. Other systems add flippers in the rear as well. Flippers eliminate the risk of an open track and enable the robot to right itself or crawl by using the flippers to push itself back over or pull itself forward.

Biomimetic Locomotion

The main concept is that there is a periodic, or repetitive, motion, either a vibration of the body or an oscillation of limbs. It can be divided into 5 main categories:

- *Crawling* occurs when the agent overcomes friction through longitudinal vibration or movement. An example in nature is a caterpillar which contracts along the length of its body to move forward, creating a slow sinusoidal wave of motion. The active Scope Camera (ASC) robot used to search the rubble at a building collapse in Jacksonville is an example of a crawling robot. Rather than duplicate the sinusoidal, the designers used vibrations to move the robot without complex mechanical deformation. The ASC is surrounded with a "skin" of VELCRO tilted at an angle similar to the "hairs" in a hairy caterpillar. As the robot vibrates throughout the axis, the VELCRO makes contact with surfaces and is thrust forward slightly due to the angle of tilt. Crawling is the most energy-intensive category of movement, so it should be no surprise that the ASC is connected to a large power supply.
- *Sliding* occurs when the agent overcomes friction through transverse vibrations or movements. An example in nature is a snake which moves in a sinusoidal side-to-side pattern using its scales to increase friction enabling pivoting. A snake robot uses transverse vibrations to move. It should be noted that most robots that are called "snakes" do not use transverse locomotion. They are actually replicating the flexibility of a snake or an elephant's trunk in order to move through tight spaces with high curvature or climb or grasp poles. The appeal of a snake or elephant trunk is that it is small but highly flexible, because it has many joints; in engineering this is called a hyper-redundant mechanism. The flexible robots often use a series of omnidirectional wheels to achieve the same mobility as a snake but without the frictional aspects.
- *Running* occurs when the agent overcomes kinetic energy with an oscillatory movement of a multi-link pendulum that leads to a predominately horizontal motion. A multi-link pendulum is an engineering representation of a leg and joints. Jumping is another way the agent can overcome kinetic energy with oscillatory movement of legs. In this case the oscillations produce a predominately vertical motion.
- *Walking*
- *Railway wheels*

Running, jumping and walking are integral to legged locomotion and form a separate field of study.

Legged Locomotion

Legs can provide versatile locomotion at the lower end energy cost for biological locomotion. In terms of AI robotics, legged locomotion has two components: computing the reference trajectory, or where the robot should go, and computing the gait (Andatura), or the specific type of oscillation to move along the reference trajectory. Computing the gait can be hard because there is a large set of possible motions for the legs, called a leg event, to create an oscillatory motion, and the oscillatory motions have to be adapted to the terrain so that the legs land on the right footfalls in order to maintain balance. Siegwart, Nourbakhsh and Scaramuzza note that the number of possible leg events, N , to achieve legged locomotion increases factorially. Given k as the number of legs, we have that $N = (2k-1)!$. We use U and D to denote if a leg moves up or down respectively.

We use - if the leg does not move. The problem of control increases if the leg has joints. Now the control signal is not just up, down or not at all but it has to signal which parts, or linkages, of the leg move and by how much. The human legged is jointed such that there are 7 DOFs, or seven possible control directives to specify, just to move the leg to a desired footfall and this ignores placement of the foot and toes.

As a result of the high number of leg events, the approach is to bundle the movements into a small set of pre-computed coordinated movements called gaits.

Balance

An obvious challenge is to maintain balance. We can create static walkers, i.e., robots that are always statically balanced. The other is provide dynamic balance.

- *Static balance*: The concept of an agent is balanced while it walks relies on the notion of a support polygon. The support polygon is defined as the convex hull of the contact points with the ground. If the center of mass of the agent is in the support polygon, then the agent is statically balanced. For bipedal locomotion, the polygon degenerates to a line, requiring more precise poses in order for the agent to remain balanced. Static stability means when a robot is statically balance when they are moving and also when they are stationary. It means that at every moment i time, the robot is in static balance. It also means that the legs maintain balance without any passive correction, or that if the robot put down a set of legs, the legs were firmly planted on the ground and would not slip. Maintaining static stability while walking is called static walking. It is typically slow since speed is subordinated to safe movements.
- *Dynamic balance* is what most of the animals rely on. Consider normal human running: there is an instant in time when neither leg is on the ground as the human has pushed off of one foot and lifted it to start swinging it forward, but the other leg has not touched down yet. A challenge in dynamic balance is to make sure that when the agent lands on the forward leg, the leg do not slip out from under causing a fall, and the legs are positioned such that the agent can spring off of them. This is commonly modeled as placing the legs so that there is zero moment point on the leg. Think of the lg as an inverted pendulum where the leg can make contact with the ground within a range of angles. The zero momentum point is the angle where the horizontal forces of momentum and friction are balanced and thus the robot should not fall. The ZMP is used to compute where to place a leg, but it is also influenced by friction and the type of foot on the leg. Just like static walking, locomotion using ZMP principle to place a leg typically has to stomp to get the foot flat against the ground to ensure that it will not slip.

Gaits

Being able to compute the ZMP is one step in dynamic balance of legged locomotion, but there is still the issue of how to reduce the number of legged events. A key concept i simplifying locomotion is organizing the oscillation of the legs into virtual gaits. The idea is that legs are collected into two sets, A and B, and all the legs in A move at the same time in the same way and then next all legs in B move in the same way at the same time. Doing this, we don't have a factorial computation but a fixed one, no matter how many legs an agent has. We have 3 basic gaits based on how legs are paired: Trot, where diagonal pairs of legs alternate moving; pace, where lateral pairs alternate moving; Bound, where front and back pairs alternate moving.

Legs with joints

Virtual gaits reduce the effort in planning the number of leg events for dynamic balance because it reduces the number of legs to two virtual sets. However, gaits generally treat legs as inverted pendulum, ignoring the fact that legs often have joints. The joint in a leg gives the agent a mechanism by which to spring up with kinetic energy and thus increase the range of motion and the ability to adapt to terrain. They also reintroduce computational complexity. Following the idea of virtual gaits, joints in a leg are generally moved according to a small set of oscillatory patterns, or macros. Whereas gaits have been studied for thousands of years in animal husbandry, the patterns for moving the joints on a single leg have not been well understood. There are two main approaches to determining the motion of legs with joints: use motion capture cameras to determine how a person or animal moves their legs for a particular situation and then duplicate that movement or develop a central pattern generator CPG for each joint.

CPG

In biology, there are neural structures called central pattern generators that produce oscillations which require synchronized movements. Once the CPG is activated, the oscillations occur without the need for sensing or other additional inputs. However, the default pattern can be modified, or tuned, with sensing. Ideally a CPG could be developed that creates the rolling polygon of leg motion up, forward, and down, and then the vertices on the polygon could be adjusted to adapt to sensed terrain. The most common implementation of a CPG is with a variant of the nonlinear Van der Pol equation. Motion capture and CPG are not mutually exclusive. We can use motion capture to document how a gait changes to adapt to a terrain which then can be incorporated as a tuning input to a CPG or a new CPG that is released as appropriate.

Action Selection

There exists a third way for locomotion: learning to assemble existing motor schemas so that an agent learns to locomote. The learning approach reinforces the principles and merit of behaviors. Action selection connotes how an agent chooses which behavior to instantiate at a given time. Innate releasing mechanisms are technically a form of action selection, but those are pre-programmed and hardwired. The bigger question is how actions are selected and learned for new events.

Locomotion from Sigwart

Nature favors legged locomotion, since locomotion systems in nature must operate on rough and unstructured terrain.

Key issues for locomotion

Locomotion is the complement of manipulation. Locomotion and manipulation thus share the same core issues of stability, contact characteristics and environmental type:

- stability: number and geometry of contact points; center of gravity; static/dynamic stability; inclination of terrain
- characteristics of contact: contact point/path size and shape; angle of contact; friction
- type of environment: structure; medium.

Legged locomotion

It is characterized by a series of point contacts between the robot and the ground. Because only a set of point contacts is required, the quality of the ground between those points does not matter as long as the robot can maintain adequate ground clearance. A walking robot is capable of crossing a hole or chasm so long as its reach exceeds the width of the hole. A final advantage is the potential to manipulate objects in the environment with great skill. The disadvantages include power and mechanical complexity. The leg, must be capable of sustaining part of the robot's total weight, and in many robots it must be capable of lifting and lowering the robot. High maneuverability will only be achieved if the legs have a sufficient number of DOFs to impart forces in a number of different directions.

The higher the number of legs, the higher the complexity of the active control to maintain balance. In contrast, a creature with three legs can exhibit a static, stable pose provided that it can ensure that its center of gravity is within the tripod of ground contact. Static stability, demonstrated by a three-legged stool, means that balance is maintained with no need for motion. In order to achieve static walking, a robot must have at least four legs, moving one of them at a time.

Compared to human legs, which need 7 DOFs, a robot leg requires a minimum of 2 DOFs to move a leg forward by lifting the leg and swinging it forward. In recent robots, they added a fourth DOF to the ankle for bipedal walking, which enables the robot to shift the resulting force vector of the ground contact by actuating the pose of the sole of the foot. In general, adding degrees of freedom to a robot leg increase the maneuverability of the robot, both augmenting the range of terrains on which it can travel and the ability of the robot to travel with a variety of gaits. The disadvantage, though, is energy, control and mass, further increasing power and load requirements on existing actuators.

Consideration of dynamics

The cost of transportation expresses how much energy a robot uses to travel a certain distance. To better compare differently sized systems, this value is usually normalized by the robot's weight and expressed in $J/(Nm)$. When a robot moves with constant speed on a level surface, its potential and kinetic energy remain constant. In theory, no physical work is necessary to keep it moving, which makes it possible to get from one place to another with zero cost of transportation. In reality, some energy is always dissipated, and robots have to be equipped with actuators and batteries to compensate for the losses. Friction is present in the joints of legged systems and energy is dissipated by the foot-ground interaction. However, these effects cannot explain why legged systems usually consume considerably more energy than their wheeled counterparts. The bulk part of energy loss actually originates in the fact that legs are not performing a continuous motion, but are periodically moving back and forth. Joints have to undergo alternating phases of acceleration and deceleration, and energy is irrecoverably lost in the process. Because of the segmented structure of the legs, it can even happen that energy that is fed into one joint is simultaneously dissipated by another joint without creating any net work at the feet. Therefore, actuators are working against each other.

A solution to this problem is a better exploitation of the dynamics of the mechanical structure. The natural oscillations of pendula and springs can automatically create the required periodic motions. During running, these inverted pendulum dynamics are additionally enhanced by springs, which store energy during the ground phase and allow the main body to take off for the subsequent flight phase. With this approach it is possible to build legged robots that do not have actuation of any kind. Such passive dynamic walkers walk down a shallow incline but, because no actuators

are present, no negative work is performed and energetic losses due to braking are eliminated. When the locomotion speed changes, characteristic properties such as stride length or stride frequency change as well, and more and more actuator effort is needed to force the joints to follow their required trajectories. Changing the gait allows us to use a different set of natural dynamics, which better matches the stride frequency and step length that are needed for higher velocities.

Wheeled mobile robots

The wheel has been by far the most popular locomotion mechanism in mobile robotics and in man-made vehicles in general. It can achieve good efficiency and it does so with relatively simple mechanical implementation. Due to the fact that wheeled robots are almost always designed so that all the wheels are in ground contact at all times, balance is not usually a research problem in the robot design. Three wheels are sufficient to guarantee stable balance but also two-wheels robot can be stable. When we have more than 3 wheels, a suspension system is required to allow all wheels to maintain ground contact when the robot encounters uneven terrain. For wheeled robots we care more about traction, stability, maneuverability and control than balance.

There are 4 major wheel classes, which differ in their kinematics. The standard wheel and the castor wheel have a primary axis of rotation and are thus highly directional. To move in different direction, the wheel must be steer first along a vertical axis. The key difference between the two is that the standard wheel can accomplish this steering with no side effects, since the center of rotation passes through the contact patch with the ground, whereas the castor wheel rotates around an offset axis, causing a force to be imparted to the robot chassis during steering.

The Swedish wheel and the spherical wheel are both designs that are less constrained by directionality than the conventional standard wheel. The Swedish wheel functions as a normal wheel but provides low resistance in another direction as well, sometimes perpendicular to the conventional direction, and sometimes at an intermediate angle. The small rollers attached around the circumference of the wheel are passive and the wheel's primary axis serves as the only actively powered joint. The key advantage of this design is that, although the wheel rotation is powered only along the one principal axis, the wheel can kinematically move with very little friction along many possible trajectories, not just forward and backwards.

The spherical wheel is a truly omnidirectional wheel, often designed so that it may be actively powered to spin along any direction. One mechanism for implementing this spherical design imitates the computer mouse, providing actively powered rollers that rest against the top surface of the sphere and impart rotational forces.

The choice of wheel types for mobile robots is strongly linked to the choice of wheel arrangement, or wheel geometry. Mobile robots are designed for applications in a wide variety of situations. The minimum number of wheels required for static stability is two. A two-wheeled differential-drive robot can achieve static stability if the center of mass is below the wheel axle. However, under ordinary circumstances, such a solution requires wheel diameters that are impractically large. Static stability requires a minimum of three wheels, with the additional caveat that the center of gravity must be contained within the triangle formed by the ground contact points of the wheels.

Some robots are omnidirectional, meaning that they can move at any time in any direction along the ground plane regardless the orientation of the robot around its vertical axis. This level of maneuverability requires Swedish or spherical wheels. In general, the ground clearance of robots with Swedish or spherical wheels is somewhat limited due to the mechanical constraints of constructing omnidirectional wheels. A interesting modern solution is the four-castor wheel configuration in which each castor wheel is actively steered and actively translated. In this configuration the robot

is truly omnidirectional.

There is generally an inverse correlation between controllability and maneuverability. For example, the omnidirectional designs such as the four-caster wheel configuration require significant processing to convert desired rotational and translational velocities to individual wheel commands. Such omnidirectional designs often have greater DOFs at the wheel. These DOFs cause an accumulation of slippage, tend to reduce dead-reckoning accuracy and increase the design complexity. Controlling an omnidirectional robot for a specific direction of travel is also more difficult and often less accurate when compared to less maneuverable designs. In summary, there is no ideal drive configuration that simultaneously maximizes stability, maneuverability and controllability. Each mobile robot application places unique constraints on the robot design problem, and the designer's task is to choose the most appropriate drive configuration possible among this space of compromises.

Synchro drive

Is a popular arrangement of wheels in indoor mobile robot applications. There are three driven and steered wheels but only two motors are used in total. The one translation motor sets the speed of all three wheels together, and the one steering motor spins all the wheels together about each of their individual steering axes. Not that the wheels are being steered with respect to the robot chassis, and therefore there is no direct way of orienting the robot chassis. In fact, the chassis orientation does drift over time due to uneven tire slippage, causing rotational dead-reckoning error.

It is particularly advantageous in cases where omnidirectionality is sought. So long as each vertical steering axis is aligned with the contact path of each tire, the robot can always orient its wheels and move along a new trajectory without changing its footprint. Of course, if the robot chassis has directionality and the designers intended to reorient the chassis purposefully, then synchro drive is appropriate only when combined with an independently rotating turret that attaches to the wheel chassis. In terms of dead reckoning, synchro drive systems are generally superior to true omnidirectional configurations but inferior to differential-drive and Ackerman steering systems (AKA cars). There are two main reasons for this: the translation motor generally drives the three wheels using a single belts. Because of slop and backlash in the drive train, whenever the drive motor engages, the closest wheel begins spinning before the furthest wheel, causing small change in the orientation of the chassis. With additional changes in motor speed, these small angular shifts accumulate to create a large error in orientation during dead reckoning. Second, the mobile robot has no direct control over the orientation of the chassis. Depending on the orientation of the chassis, the wheel thrust can be highly asymmetric, with two wheels on one side and the third wheel alone. or symmetric, with one wheel on each side and one wheel straight ahead or behind. The asymmetric case result in a variety of errors when tire-ground slippage can occur, again causing errors in dead reckoning of robot orientation.

Omnidirectional drive

Omnidirectional robots are able to move in any direction (x, y, θ) at any time are also holonomic. They can be realized by using spherical, caster or Swedish wheels:

- **3 spherical wheels:** each wheel is actuated by one motor. In the design presented by the book, the spherical wheels are suspended by three contact points, two given by spherical bearings and one by a wheel connected to the motor axle. This concept provides excellent maneuverability and is simple in design. However, it is limited to flat surfaces and small loads, and it is quite difficult to find round wheels with high friction coefficients.

- **4 Swedish wheels(Carnegie Mellon Uranus):** The configuration in the book uses 4 Swedish 45-degree wheels, each driven by a separate motor. By varying the direction of the rotation and relative speed of the four wheels, the robot can be moved along any trajectory in the plane and can simultaneously spin around its vertical axis. This arrangement is not minimal in terms of control motors. Because there are only 3 DOFs in the plane, one can build a 3 wheel omnidirectional robot chassis using 3 Swedish 90-degrees wheels.
- **4 castor wheels and 8 motors(Nomad XR4000)**

Tracked slip/skid locomotion

In the previous configurations we have made the assumption that wheels are not allowed to skid against the surface. An alternative form of steering, termed slip/skid, may be used to reorient the robot by spinning wheels that are facing the same direction at different speeds or in opposite directions. Robots that make use of track have much larger ground contact patches, and this can be significantly improve their maneuverability in loose terrain compared to conventional wheel designs. However, due to this large ground contact patch, changing the orientation of the robot usually requires a skidding turn, wherein a large portion of the track must slide against the terrain. The disadvantage of such configurations is coupled to the slip/skid steering. Because of the large amount of skidding during a turn, the exact center of rotation of the robot is hard to predict and the exact change in position and orientation is also subject to variations depending on the ground friction. Therefore, dead-reckoning on such robots is highly inaccurate. Furthermore, a slip/skid approach on a high-friction surface can quickly overcome the torque capabilities of the motors being used. In terms of power efficiency, this approach is reasonably efficient on loose terrain but extremely inefficient otherwise.

Walking wheels

Walking robots might offer the best maneuverability in rough terrain. However, they are inefficient on flat grounds and need sophisticated control. Hybrid solutions combining the adaptability of legs with the efficiency of the wheels, offer an interesting compromise.

Kinematics

Kinematics from notes

trycicles

The circular path of a car-like tricycle robot is described by the following equations

$$R = \frac{L}{\tan s}$$

$$R_d = \frac{L}{\sin s}$$

Where R_d is the radius of the path that the rear driving wheel moves, while R is the distance from the robot center and the center of the circular path described by the robot motion. While s is the steering angle, i.e., the angle between the robot main axis and the axis of the wheel. L is the distance of the steering wheel from the center of the robot. By changing the steering angle, we change the radius of the circular path. If s is 0, the robot moves in a straight line. If we think about a time Δt , the angle spanned by the robot $\Delta\theta$ is the following:

$$\Delta\theta = \frac{v\Delta t}{R_d} = \frac{v\Delta t \sin s}{L}$$

We obtain the following two equations.

$$R = \frac{L}{\tan s}, \Delta\theta = \frac{v\Delta t \sin s}{L}$$

by controlling the velocity and the steering angle we can control the motion of the robot on the plane.

In real life, we use electric motors, but they have low torque and high speed. So to augment the torque we use gears of different size. The point of doing so is that by applying the same speed to a different gear we can increase the torque generated by the electric motor. The relation that governs the gears is the following:

$$\omega_2 = \frac{r_1}{r_2} \omega_1$$

It derives from the fact that the two gears have the same speed. When a small gear drives a bigger gear, the second gear has an higher torque and lower angular velocity in proportion to the ratio of teeth. Gears can be chained together to achieve compound effect. The angular speed is not the only changing factor, also the torques changes with this mechanism:

$$t_2 = \frac{r_2}{r_1} t_1$$

which has the reciprocal constant factor compared to the angular speed equation.

The robot we deal with has no static reference from which we can derive everything as in the robot manipulator case. We talk about forward and inverse kinematics in the sense that by knowing some information about the state of the robot we can derive the final positioning of the robot. For a mobile robot the position of the motors does not tell the whole story, because there is no direct way to measure the robot position in the environment instantaneously. We can integrate position overtime, we can estimate the robot pose by looking at how the wheels moved through time in order to find the position of the robot, but this is prone to inaccurate measurements, which accumulates and leads to a drift (not sure about this last word). Understanding mobile robot motion starts from understanding wheel's constraints placed on robot mobility. So, a manipulator, by having a fixed position compared to the global frame, in order to calculate the position of the end-effector is enough to know the structure of the robot and the position of the joints, which can be calculated by encoders placed inside the joints. In the case of our robot, it is not enough to have encoders in the wheels to determine the final position of the robot since we don't have a fixed transformation in the reference system. With non-holonomic systems it is not possible to calculate the final position by looking at the position of the encoders, but we need to look at the differential equations (aka speeds) to find the final position of the robot. The traveled distance in the space is not enough to determine the position of the robot.

We define forward kinematic as a transform from the joint space to the physical space while the inverse kinematic is a transformation from the physical space to the joint space. The latest is the one required for motion control cause most of the time I want to bring my robot to a certain configuration on the plane, so I want to actuate the motors to reach this position. Since our systems are subject to non-holonomic constraint we need to use inverse kinematics (differential kinematics) so it means that we need to deal with velocities, so how the position changes in time, and not the absolute change in position. So, to control our robot we use the block that is in SLIDE 12, LOOK AT THE PICTURE BITCH. Since the robot does have mass, it cannot change its speeds from 0 to v or viceversa, so the inverse kinematics does not take into account the mass of the robot itself (we assume that the robot has no mass, it's a cheat code or so). So, inverse kinematics does not take into account forces and accelerations. So dynamics is another story. We represent the robot in the world reference frame and in the robot reference frame. We want to be able to express the robot position as $\psi = [\dot{x} \ \dot{y} \ \dot{\theta}]^T$, so we would like to have a mapping between the two frames $\dot{\psi}_R = R(\theta)\dot{\psi}_W$, where R is the rotational matrix (see Segwart notes). CASE OF ROBOT ALIGNED WITH Y, LOOK AT GIANNINI'S NOTES.

The goal of the kinematic model is to be able to determine the robot speed as a function of the speeds of the wheels, the angle of the steering (if there), and the speed at which we steer the wheel. We want to describe, more generally, the robot speed $\dot{\psi}$ as a function of wheel speed $\dot{\phi}_i$, where i notes the wheel, steering angles β_i , and steering speeds $\dot{\beta}_i$. All these variables are connected by the geometrical parameters of the robot chassis (the constants of our model).

When we speak of the forward kinematics we need the following:

$$\dot{\psi} = [\dot{x} \ \dot{y} \ \dot{\theta}] = f(\dot{\phi}_1, \dots, \dot{\phi}_n, \beta_1, \dots, \beta_m, \dot{\beta}_1, \dots, \dot{\beta}_m)$$

The function f describes the geometrical meaning of the robot. The inverse kinematic instead is described as follows:

$$[\dot{\phi}_1, \dots, \dot{\phi}_n, \beta_1, \dots, \beta_m, \dot{\beta}_1, \dots, \dot{\beta}_m] = g(\dot{x}, \dot{y}, \dot{\theta})$$

How can we write all this in the simple case of the differential drive robot? given l as the semi-distance between the wheels and r as the radius of the wheel and θ is the heading of the robot., we can write some fancy equations that are on the slide, but I'm too lazy to write (SLIDE 3).

Degree of maneuverability

The degree of maneuverability it's the motion capabilities that our platform has. There are the motion variables that we can control in order to move our robot. To decide the next pose of our robot. In other words, they are the overall DOFs that the robot can manipulate. It is obtained by summing the degree of mobility and degree of steerability. The degree of mobility it's δ_m , while the degree of steerability it's δ_s . So $\delta_M = \delta_m + \delta_s$. The degree of mobility are the DOFs of the chassis that can be control by changing the speed of the wheels. The degrees of steerability are the DOFs allowed by the change in orientation of the steerable wheels. The latter is the particular change at any instant of a steerable wheel imposes a (different) kinematic constraint.

Some examples are:

- *Unicycle*: has 1 single fixed standard wheel. $\delta_s = 0, \delta_m = 1$, so $\delta_M = 1$
- *Differential drive*: has 2 fixed standard wheels with two different arrangements (wheels on different axis or wheels on the same axis). $\delta_s = 0, \delta_m = 2$, so $\delta_M = 2$ in the first case, $\delta_s = 0, \delta_m = 1$, so $\delta_M = 1$ in the second case.
- *Omnidirectional robot*: with 3 Swedish wheels, each one of them driven by a different motor, here $\delta_m = 3$
- *bicycle*: 1 fixed standard wheel and 1 steerable standard wheel has $\delta_s = 1, \delta_m = 1, \delta_M = 2$
- *Tricycle*: 1 steerable wheel and 2 fixed standard wheel has $\delta_s = 1, \delta_m = 2$ (if the two fixed are connected by a rod)
- *two-steer vehicle*: 2 steerable wheels and 1 spherical wheel for support has $\delta_s = 2, \delta_m = 1$

Two robots with the same degree of maneuverability δ_M are not equal nor move in the same way. For any robot with $\delta_M = 2$ the ICR is always constrained to lie on a line. for any robot with $\delta_M = 3$, the ICR is not constrained and can be set anywhere on the plane. Castor wheels do not add any constraint in steerability. They are useless.

We have five basic types of 3 wheel configurations: (PICTURE on the slides)

- *Omnidirectional drive*: we have 3 Swedish wheels $\delta_s = 0, \delta_m = 3, \delta_M = 3$
- *Differential drive*: $\delta_s = 0, \delta_m = 2, \delta_M = 2$
- *Omnisteer*: 2 Swedish wheel plus a steerable $\delta_s = 1, \delta_m = 2, \delta_M = 3$
- *tricycle*: $\delta_s = 1, \delta_m = 1, \delta_M = 2$
- *two-steer drive*: two steerable wheel with one motors and one steerable wheel as support. It has $\delta_s = 2, \delta_m = 1, \delta_M = 3$

The synchro drive has one motor to steer the steering belt and one robot to power the wheels, so $\delta_m = 1, \delta_s = 1, \delta_M = 2$. It has 2 DOF in the environment (x,y) and not θ . The ICR it's always at infinity. In fact this robot it's never rotating but always traslating. This is another reason why this robot can never change its heading.

The DOFs of the motion of a robot is the robot's ability to achieve various poses in the environment or in its working space. (Pose its position plus orientation). For example the car has $\delta_M = 2$ but the car has 3 DOFs in the environment. The environment DOFs tells you how much constrained is the degrees of mobility of your robot.

$\delta_m \Rightarrow$ robot's independently achievable velocities, which is the number of dimensions of the velocity space, which is the differentiable degrees of freedom DDOF. If DDOF=3 the robot can manipulate the 3 variable for the pose(x,y,theta with the dot on top).

DOFs are the robot ability to achieve various poses. DDOFs are the robot's ability to achieve various velocities in the velocity space,i.e., the ability to achieve various trajectories. The rule is $DDOF \leq \delta_m \leq DOF$. Only robot with zero non-holonomic constraints are omnidirectional robot, and they have DDOF=DOF.

Mobile robot workspace

We need to introduce some concepts first:

- **Position:** the location of the robot in the 3D space; for example in the 2D space, with a mobile robot on a plane, you would use x and y to determine the position of the robot. In the case of a flying robot you use x,y and z;
- **Pose:** It captures both the position and the orientation of the robot. It is the value of the state variables of the robot in the space of its DOFs. In the case of a mobile robot on the plane we have that the pose is represented by x, y, θ . In the case of a flying robot we would have 6 variables to determine the pose.
- **Path:** It's the sequence of poses in the space of the robot's DOFs.
- **Trajectory:** when the robot was in a specific position. It adds time to the path. It is the path + the time. You add another variable, which is the time stamp at which the robot was in that position.

The robots can have same paths but different trajectories.

Open loop control

A kinematic controller can compute the motion commands for a robot platform to follow a given trajectory described by poses vs. time. The trajectory is divided i motion segments of clearly defined shape: straight lines and segments of a circle. The control problem is. to pre-compute a smooth trajectory based on line and circle segments using the kinematics of the mobile platform. The disadvantages: are the following:

- it is not at all an easy task to pre-compute a feasible trajectory
- limitations and constraints of the robot's velocities and accelerations
- does not adapt or correct the trajectory if dynamical changes in the environment occur
- the resulting trajectories are usually not smooth
- the real position of the robot is not fed back to the controller.

Closed loop control

A feedback controller is the solution to control a mobile platform to move from point A to point B. The robot's path-planning is reduced to setting intermediate positions laying on the requested path. We need to find a control matrix K , if exists:

$$K = \begin{bmatrix} k_{11} & k_{12} & k_{13} \\ k_{21} & k_{22} & k_{23} \end{bmatrix} \quad k_{ij} = k(t, e)$$

such that the control of $v(t)$ and $\omega(t)$ brings the error to zero: $e(t)$.

$$\begin{bmatrix} v(t) \\ \omega(t) \end{bmatrix} = Ke = K \begin{bmatrix} x \\ y \\ \theta \end{bmatrix}$$

$$\lim_{t \rightarrow \infty} e(t) = 0$$

To move the robot from one place to the other we need the Trajectory following (and obstacle Avoidance), using sensors and robot's current positions as inputs of the module. Before following the trajectory we need a higher level module to plan a trajectory, the Trajectory Planning module, which takes as input the map of the environment and the goal position, and the current position. The trajectory planning module has a lower frequency, while the trajectory following module has a higher frequency. On the slides there is a cute image of the trajectory following control loop. To lazy to add it now.

Kinematics from Segwart

We need a clear mapping between global and local reference frames, since the forces and constraints of each wheel must be expressed with respect to a clear and consistent reference frame.

Throughout the analysis we model the robot as a rigid body on wheels, operating on a horizontal plane. The total dimensionality of this robot chassis on the plane is 3, two for position in the plane and one for orientation along the vertical axis, which is orthogonal to the plane. Of course, there are additional DOFs and flexibility due to the wheel axles, wheel steering joints and wheel castor joints. However, by robot chassis we refer only to the rigid body of the robot, ignoring the joints and DOFs internal to the robot and its wheels.

In order to specify the position of the robot on the plane we establish a relationship between the global reference frame of the plane and the local reference frame of the robot. The axes X_I, Y_I define an arbitrarily inertial basis on the plane as the global reference frame from some origin $O : \{X_I, Y_I\}$. To specify the position of the robot, choose a point P on the robot chassis as its position reference point. The basis $\{X_R, Y_R\}$ defines two axes relative to P on the robot chassis and is thus the robot's local reference frame. The position of P in the global reference frame is specified by coordinates x and y , and the angular difference between the global and local reference frames is given by θ . We can describe the pose of the robot as a vector with these elements. Note the use of the subscript I to clarify the basis of this pose as the global reference frame:

$$\xi_I = \begin{bmatrix} x \\ y \\ \theta \end{bmatrix}$$

To describe robot motion in terms of component motions, it will be necessary to map motion along the axes of the global reference frame to motion along the axes of the robot's local reference frame.

Of course, the mapping function is a function of the current pose of the robot. This mapping is accomplished using the orthogonal rotation matrix:

$$R(\theta) = \begin{bmatrix} \cos\theta & \sin\theta & 0 \\ -\sin\theta & \cos\theta & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

This matrix can be used to map motion in the global reference frame $\{X_i, Y_i\}$ to motion in terms of the local reference frame $\{X_R, Y_R\}$. This operation is denoted by $R(\theta)\xi_I$ because the computation of this operation depends on the value of θ :

$$\dot{\xi}_R = R\left(\frac{\pi}{2}\right)\dot{\xi}_I$$

Forward kinematics model

In the simplest case, the mapping showed previously is sufficient to generate a formula that captures the forward kinematics of the mobile robot.

Now, consider a differential drive robot with two wheels, each with diameter r . Given a point P centered between the two drive wheels, each wheel is a distance l from P . Given r, l, θ , and the spinning speed of each wheel $\dot{\phi}_1$ and $\dot{\phi}_2$, a forward kinematic model would predict the robot's overall speed in the global reference frame:

$$\dot{\xi}_I = \begin{bmatrix} \dot{x} \\ \dot{y} \\ \dot{\theta} \end{bmatrix} = f(l, r, \theta, \dot{\phi}_1, \dot{\phi}_2)$$

We know that we can compute the robot's motion in the global reference frame from motion in its local reference frame: $\dot{\xi}_I = R(\theta)^{-1}\dot{\xi}_R$. Therefore, the strategy will be to first compute the contribution of each of the two wheels in the local reference frame, $\dot{\xi}_R$. For this example of a differential-drive chassis, this problem is particularly straightforward.

Suppose that the robot's local reference frame is aligned such that the robot moves forward along $+X_R$. First consider the contribution of each wheel's spinning speed to the translation speed at P in the direction of $+X_R$. If one of the wheel spins while the other wheel contributes nothing and is stationary, since P is halfway between the two wheels, it will move instantaneously with half the speed: $\dot{x}_{r_1} = (1/2)r\dot{\phi}_1$ and $\dot{x}_{r_2} = (1/2)r\dot{\phi}_2$. In a differential drive robot, these two contributions can simply be added to calculate \dot{x}_R component of $\dot{\xi}_R$.

Consider now a differential drive robot in which each wheel spins with equal speed but in opposite direction. The result is a stationary spinning robot. As expected, \dot{x}_R will be zero as \dot{y}_R . Finally, we must compute the rotational component $\dot{\theta}_R$ of $\dot{\xi}_R$. Once again, the contributions of each wheel can be computed independently and just adds. Consider the right wheel (wheel 1.) Forward spin of this wheel results in counterclockwise rotation at point P . The rotational velocity ω_1 at P can be computed because the wheel is instantaneously moving along the arc of a circle of radius $2l$: $\omega_1 = \frac{r\dot{\phi}_1}{2l}$. The same calculation applies to wheel 2, with the exception that forward spin results in clockwise rotation at point P : $\omega_2 = \frac{-r\dot{\phi}_2}{2l}$. Combining these two individual formulas yields a kinematic model for the differential-drive example robot:

$$\dot{\xi}_I = R(\theta)^{-1} \begin{bmatrix} \frac{r\dot{\phi}_1}{2} + \frac{r\dot{\phi}_2}{2} \\ 0 \\ \frac{r\dot{\phi}_1}{2l} + \frac{-r\dot{\phi}_2}{2l} \end{bmatrix}$$

We can now use this kinematic model in an example. However, we must first compute $R(\theta)^{-1}$. In general, calculating the inverse of a matrix may be challenging. In this case, however, it is easy because it is simply a transform from $\dot{\xi}_R$ to $\dot{\xi}_I$ rather than vice versa:

$$R(\theta)^{-1} = \begin{bmatrix} \cos\theta & -\sin\theta & 0 \\ \sin\theta & \cos\theta & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Wheel kinematic constraints

We assume that the plane of the wheel always remains vertical and that there is in all cases one single point of contact between the wheel and the ground plane. We assume that there is no sliding at this single point of contact. That is, the wheel undergoes only under conditions of pure rolling and rotation about the vertical axis through the contact point. We have two constraints for every wheel type: the first one is that the wheel must roll when motion takes place in the appropriate direction, the second is that the wheel must not slide orthogonal on the wheel plane.

fixed standard wheel

It has no vertical axis of rotation for steering, so its angle to the chassis is thus fixed, and it is limited to motion back and forth along the wheel plane and rotation around its contact point with the ground plane. Let us consider the following case, where the fixed standard wheel A is expressed in polar coordinates by distance l and angle α . The angle of the wheel plane relative to the chassis is denote by β , which is fixed since the standard wheel is not steerable. The wheel, which has radius r , can spin over time, and so its rotational position around its horizontal axle is a function of time t : $\phi(t)$. The rolling constraint for this wheel enforces that all motion along the direction of the wheel plane must be accompanied by the appropriate amount of wheel spin so that there is pure rolling at the contact point:

$$[\sin(\alpha + \beta) - \cos(\alpha + \beta)(-l)\cos\beta]R(\theta)\dot{\xi}_I - r\dot{\phi} = 0$$

The first term of the sum denotes the total motion along the wheel plane. The three elements of the vector on the left represent mapping from each $\dot{x}, \dot{y}, \dot{\theta}$ to their contributions for motion along the wheel plane. The sliding constraint for this wheel enforces that the component of the wheel's motion orthogonal to the wheel plane must be zero:

$$[\cos(\alpha + \beta)\sin(\alpha + \beta)l\sin\beta]R(\theta)\dot{\xi}_I = 0$$

Steered standard wheel

They differ from the fixed standard wheel only in that there is an additional degree of freedom: the wheel may rotate around a vertical axis passing through the center of the wheel and the ground contact point. The equations of position for the steered standard wheel are identical to that of the fixed standard wheel with one exception: the orientation of the wheel to the robot chassis is no longer a single fixed value β , but instead varies as a function of time: $\beta(t)$. The rolling and sliding constraints are:

$$[\sin(\alpha + \beta) - \cos(\alpha - \beta)(-l)\cos\beta]R(\theta)\dot{\xi}_I - r\dot{\phi} = 0$$

$$[\cos(\alpha + \beta)\sin(\alpha + \beta)l\sin\beta]R(\theta)\dot{\xi}_I = 0$$

These constraints are identical to those of the fixed standard wheel because, unlike $\dot{\phi}$, $\dot{\beta}$ does not have a direct impact on the instantaneous motion constraints of a robot.

Castor wheels

They are able to steer around a vertical axis, however, unlike the steered standard wheel, the vertical axis of rotation in a castor wheel does not pass through the ground contact point. The wheel contact point is now at position B, which is connected by a rigid rod AB of fixed length d to point A fixes the location of the vertical axis about which B steers, and this point A has a position specified in the robot's reference framed. We assume that the plane of the wheel is aligned with AB at all times. The castor wheel has two parameters that vary as a function of time. $\phi(t)$ represents the wheel spin over time as before. $\beta(t)$ denotes the steering angle and orientation of AB over time. The rolling constraint is identical to the one of the steered standard wheel since the offset axis plays no roll during motion that is aligned with the wheel plane:

$$[\sin(\alpha + \beta) - \cos(\alpha + \beta)(-l)\cos\beta]R(\theta)\dot{\xi}_I - r\dot{\phi} = 0$$

The castor geometry does have significant impact on the sliding constraint. The critical issue is that the lateral force on the wheel occurs at point A because this is the attachment point of the wheel to the chassis. Because of the offset ground contact point relative to A, the constraint that there be zero lateral movement would be wrong. Instead, the constraint is much like a rolling constraint, in that appropriate rotation of the vertical axis must take place:

$$[\cos(\alpha + \beta)\sin(\alpha + \beta)d + l\sin\beta]R(\theta)\dot{\xi}_I + d\dot{\beta} = 0$$

In this last constraint, any motion orthogonal to the wheel plane must be balanced by an equivalent and opposite amount of castor steering motion. This result is critical to the success of castor wheels because by setting the value of $\dot{\beta}$ any arbitrary lateral motion can be acceptable. In a steered standard wheel, the steering action does not by itself cause a movement of the robot chassis. But in a castor wheel the steering action itself moves the robot chassis because of the offset between the ground contact point and the vertical axis of rotation. More concisely, given any robot chassis motion $\dot{\xi}_I$, there exists some value for spin speed $\dot{\phi}$ and steering speed $\dot{\beta}$ such that the constraints are met. Therefore, a robot with only castor wheels can move with any velocity in the space of possible robot motions. We term such systems *omnidirectional*.

Swedish wheels

They have a vertical axis of rotation but they are still able to move omnidirectionally like castor wheels. Swedish wheels consist of a fixed standard wheel with rollers attached to the wheel perimeter with axes that are antiparallel to the main axis of the fixed wheel component. The exact angle γ between the roller axes and the main axis can vary. The pose of a Swedish wheel is expressed exactly as in a fixed standard wheel with the addition of a term γ representing the angle between the main wheel plane and the axis of rotation of the small circumferential rollers. The motion constraint is the following:

$$[\sin(\alpha + \beta + \gamma) - \cos(\alpha + \beta + \gamma)(-l)\cos(\beta + \gamma)]R(\theta)\dot{\xi}_I - r\dot{\phi}\cos\gamma = 0$$

Orthogonal to this direction the motion is not constrained because of the free rotation $\dot{\phi}_{sw}$ of the small rollers.

$$[\cos(\alpha + \beta + \gamma)\sin(\alpha + \beta + \gamma)l\cos(\beta + \gamma)]R(\theta)\dot{\xi}_I - r\dot{\phi}\sin\gamma - r_{sw}\dot{\phi}_{sw} = 0$$

The behavior of this constraint changes dramatically as the value γ varies.

Spherical wheels

It do not place direct constraints on motion. Such a mechanism has no principal axis of rotation, and therefore no appropriate rolling or sliding constraints exist. As with castor wheels and Swedish wheels, the spherical wheel is clearly omnidirectional and places no constraints on the robot chassis kinematics. Therefore, the following equation simply describes the roll rate of the ball in the direction of motion v_A of point A of the robot:

$$[\sin(\alpha + \beta) - \cos(\alpha + \beta)(-l)\cos\beta]R(\theta)\dot{\xi}_I - r\dot{\phi} = 0$$

By definition the wheel rotation orthogonal to this direction is zero.

$$[\cos(\alpha + \beta)\sin(\alpha + \beta)l\sin\beta]R(\theta)\dot{\xi}_I = 0$$

This equations are the same as the one for the standard fixed wheel.

Robot kinematic constraints

Given a mobile robot with M wheels we can now compute the kinematic constraints of the robot chassis. The key idea is that each wheel imposes zero or more constraints on the robot motion, and so the process is simply one of appropriately combining all of the kinematic constraints arising from all of the wheels based on the placement of whose wheels on the robot chassis.

Only fixed standard wheels and steerable standard wheels have impact on robot chassis kinematics and therefore require consideration when computing the robot's kinematic constraints. Suppose that the robot has a total of N standard wheels, comprising N_f fixed standard wheels and N_s steerable standard wheels. We use $\beta_s(t)$ to denote the variable steering angles of the N_s , while β_f to refer to the orientation of the N_f fixed standard wheels. In the case of wheel spin, both the fixed and steerable wheels have rotational positions around the horizontal axle that vary as a function of time. We denote the fixed and steerable cases separately as $\phi_f(t)$ and $\phi_s(t)$, and use $\phi(t)$ as an aggregate matrix that combines both values: $\phi(t) = \begin{bmatrix} \phi_f(t) \\ \phi_s(t) \end{bmatrix}$. The rolling constraints of all wheels can now be collected in a single expression:

$$J_1(\beta_s)R(\theta)\dot{\xi}_I - J_2\dot{\phi} = 0$$

J_2 is a constant diagonal NxN matrix whose entries are radii r of all standard wheels. $J_1(\beta_s)$ denotes a matrix with projections for all wheels to their motions along their individual wheel planes: $J_1(\beta_s) = \begin{bmatrix} J_{1f} \\ J_{1s}(\beta_s) \end{bmatrix}$. Note that $J_1(\beta_s)$ is only a function of β_s and not β_f . This is because the orientations of steerable standard wheels vary as a function of time, whereas the orientations of fixed standard wheels are constant. J_{1f} is therefore a constant matrix of projections for all fixed standard wheels.

We use the same technique to collect the sliding constraints of all standard wheels into a single expression:

$$C_1(\beta_s)R(\theta)\dot{\xi}_I = 0 \quad (3.26)$$

$$C_1(\beta_s) = \begin{bmatrix} C_{1f} \\ C_{1s}(\beta_s) \end{bmatrix}$$

This sliding constraint over all standard wheels has the most significant impact on defining the overall maneuverability of the robot chassis.

Maneuverability

The kinematic mobility of a robot chassis is its ability to directly move in the environment. The basic constraint limiting mobility is the rule that every wheel must satisfy its sliding constraint. In addition to instantaneous kinematic motion, a mobile robot is able to further manipulate its position by steering steerable wheels.

Equation 3.26 imposes the constraint that every wheel must avoid any lateral slip. Of course this holds separately for each and every wheel, and so it is possible to specify this constraint separately for fixed and for steerable standard wheels:

$$C_{1f}R(\theta)\dot{\xi}_I = 0$$

$$C_{1s}(\beta_s)R(\theta)\dot{\xi}_I = 0$$

For both of these constraints to be satisfied, the motion vector $R(\theta)\dot{\xi}_I$ must belong to the null space of the projection matrix $C_1(\beta_s)$, which is simply a combination of C_{1f} and C_{1s} . Mathematically, the null space of $C_1(\beta_s)$ is the space N such that for any vector n in N, $C_1(\beta_s)n = 0$. If the kinematic constraints are to be honored, then the motion of the robot must always be within this space N. Consider a single standard wheel. It is forced by the sliding constraint to have zero lateral motion. This can be shown geometrically by drawing a zero motion line through its horizontal axis, perpendicular to the wheel plane. At any given instant, wheel motion along the zero motion line must be zero. In other words, the wheel must be moving instantaneously along some circle of radius R such that the center of that circle is located on the zero motion line. This center point, called the instantaneous center of rotation, may lie anywhere along the zero motion line. When R is at infinity, the wheel moves in a straight line.

Robot chassis kinematics is a function of the set of independent constraints arising from all standard wheels. The mathematical interpretation of independence is related to the rank of a matrix. Recall that the rank of a matrix is the smallest number of independent rows or columns. Equation 3.26 represents all sliding constraints imposed by the wheels of the mobile robot. Therefore $rank[C_1(\beta_s)]$ is the number of independent constraints. The greater the number of independent constraints, the more constrained is the mobility of the robot. For example, in the case of a robot with a single fixed standard wheel we'll have the following relation:

$$C_1(\beta_s) = C_{1f} = [\cos(\alpha + \beta)\sin(\alpha + \beta)l\sin\beta]$$

Now let us add additional fixed standard wheel to create a differential-drive robot by constraining the second wheel to be aligned with the same horizontal axis as the original wheel. Without loss of generality, we can place point P at the midpoint between the centers of the two wheels. Given α_1, β_1, l_1 for wheel w_1 and α_2, β_2, l_2 for wheel w_2 , it holds geometrically that $\{l_1 = l_2, (\beta_1 = \beta_2 = 0), (\alpha_1 + \pi = \alpha_2)\}$. Therefore matrix $C_1(\beta_s)$ has two constraints but a rank of one:

$$C_1(\beta_s) = C_{1f} = \begin{bmatrix} \cos(\alpha_1) & \sin(\alpha_1) & 0 \\ \cos(\alpha_1 + \pi) & \sin(\alpha_1 + \pi) & 0 \end{bmatrix}$$

In general, if $rank[C_{1f}] > 1$ then the vehicle can only travel along a circle or along a straight line. This configuration means that the robot has two or more independent constraints due to fixed standard wheels that do not share the same horizontal axis of rotation.

In general, a robot will have zero or more fixed standard wheels and zero or more steerable standard wheels. We can therefore identify the possible range of rank values for any robot: $0 \leq rank[C_1(\beta_s)] \leq 3$. In case the rank is 0, there are neither fixed nor steerable standard wheels

attached to the robot frame: $N_f = N_s = 0$. In case the rank is 3, the kinematic constraints are specified along 3 DOFs, therefore there cannot be more than 3 independent constraints since the robot is completely constrained in all directions and is degenerate since motion in the plane is totally impossible. Now we can formally define a robot's *degree of mobility* δ_m :

$$\delta_m = \dim N[C_1(\beta_s)] = 3 - \text{rank}[C_1(\beta_s)]$$

The dimensionality of the null space of matrix $C_1(\beta_s)$ is a measure of the number of DOFs of the robot chassis that can be immediately manipulated through changes in wheel velocity. It is logical that δ_m must range between 0 and 3.

Degree of steerability

The degree of mobility defined above quantifies the degrees of controllable freedom based on changes to wheel velocity. Steering can also have an eventual impact on a robot chassis pose ξ , although the impact is indirect because after changing the angle of a steerable standard wheel, the robot must move for the change in steering angle to have impact on the pose. As with mobility, we care about the number of independently controllable steering parameters when defining the *degree of steerability* δ_s :

$$\delta_s = \text{rank}[C_{1s}(\beta_s)]$$

Recall that in the case of mobility, an increase in the rank implied more kinematic constraints and thus a less mobile system. In the case of steerability, an increase in the rank implies more degrees of steering freedom and thus greater eventual maneuverability. Since $C_1(\beta_s)$ includes $C_{1s}(\beta_s)$, this means that a steered standard wheel can both decrease mobility and increase steerability. The range of δ_s can be specified: $0 \leq \delta_s \leq 2$. If it is 0, then the robot has no steerable standard wheels and $N_s = 0$. The case where it is 1 is most common when a robot configuration includes one or more steerable standard wheels. In the case it is 2 occurs only when the robot has no fixed standard wheels: $N_f = 0$.

Robot maneuverability

The overall DOFs that a robot can manipulate, called the *degree of maneuverability* δ_M can be readily defined in terms of mobility and steerability:

$$\delta_M = \delta_m + \delta_s$$

Therefore maneuverability includes both the DOFs that the robot manipulates directly through wheel velocity and the DOFs that it indirectly manipulates by changing the steering configuration and moving. Note that two robots with the same δ_M are not necessarily equivalent. For example differential drive and tricycle geometry have equal $\delta_M = 2$. In the first case we have $\delta_m = 2$ and $\delta_s = 0$, while in the latter we have $\delta_m = 1$ and $\delta_s = 1$. For a robot with $\delta_M = 2$ the inner circle radius is always constrained to lie on a line, while for any robot with $\delta_M = 3$ it can be set on any point on the plane.

Mobile robot workspace

For a robot, maneuverability is equivalent to its control degrees of freedom. But the robot is situated in some environment, and the next question is to situate our analysis in the environment. Identifying a robot's space of possible configurations is important because surprisingly it can exceed δ_M . In addition to workspace, we care about how the robot is able to move between various configurations.

Degrees of freedom

In defining the workspace of a robot, it is useful to first examine its admissible velocity space. Give the kinematic constraints of the robot, its velocity space describes the independent components of robot motion that the robot can control. The number of dimensions in the velocity space of a robot is the number of independently achievable velocities. This is also called the *differentiable degrees of freedom*. A robot's DDOF is always equal to its degree of mobility δ_m . Clearly there is an inequality relation at work: $DDOF \leq \delta_m \leq DOF$.

Holonomic robots

The term holonomic in robotics, refers specifically to the kinematic constraints of the robot chassis. A holonomic robot is a robot that has zero nonholonomic constraints of the robot chassis. Conversely, a nonholonomic robot is a robot with one or more nonholonomic kinematic constraints. A holonomic kinematic constraint can be expressed as an explicit function of position variables only. A nonholonomic kinematic constraint requires a differential relationship, such as the derivative of the position variables only. Because of this latter pov, nonholonomic systems are often called nonintegrable systems. Consider the fixed standard wheel sliding constraint:

$$[\cos(\alpha + \beta) \sin(\alpha + \beta) l \sin \beta] R(\theta) \dot{\xi}_I = 0$$

This constraint must use robot motion $\dot{\xi}$ rather than pose ξ because the point is to constraint robot motion perpendicular to the wheel plane to be zero. The constraint is not integrable, depending explicitly on robot motion. Therefore, the sliding constraint is a nonholonomic constraint. The first type of holonomic robot is a robot where constraints do exist but are all holonomic kinematic constraints $\delta_m < 3$. The second type of holonomic robot exists when there are no kinematic constraint, that is $N_f = 0$ and $N_s = 0$. Since there are no kinematic constraints, there are also no nonholonomic kinematic constraints and so such robot is always holonomic.

An alternative way to describe holonomic robot is based on the relationship between the differential DOFs of a robot and the DOFs of its workspace: a robot is holonomic if and only if $DDOF = DOF$.

kinematic control

The objective of a kinematic controller is to follow a trajectory described by its position or velocity profile as a function of time. This is often done by dividing the trajectory in motion segments of clearly defined shape. The control problem is thus to precompute a smooth trajectory based on line and circle segments which drives the robot from the initial position to the final position. This approach can be regarded as open-loop motion control, because the measured robot position is not fed back for velocity or position control. It has several disadvantages:

- it is not at all an easy task to precompute a feasible trajectory if all limitations and constraints of the robot's velocities and accelerations have to be considered.
- The robot will not automatically adapt or correct the trajectory if dynamic changes of environment occur.
- The resulting trajectories are usually not smooth, because the transitions from one trajectory segment to another are, for most of the commonly used segments, not smooth. This means there is a discontinuity in the robot's acceleration.

Feedback control

A more appropriate in motion control of a mobile robot is reduced to setting intermediate positions lying on the requested path. Consider an arbitrary position and orientation of the robot and a predefined goal position and orientation. The actual pose error vector given in the robot reference frame $\{X_R, Y_R, \theta\}$ is $e = {}^R [x, y, \theta]^T$ with x, y, θ being the goal coordinates of the robot. The task of the controller layout is to find a control matrix K if it exists:

$$K = \begin{bmatrix} k_{11} & k_{12} & k_{13} \\ k_{21} & k_{22} & k_{23} \end{bmatrix} \text{ with } k_{ij} = k(t, e)$$

such that the control of $v(t)$ and $\omega(t)$ can be expressed $\begin{bmatrix} v(t) \\ \omega(t) \end{bmatrix} = Ke = K^R \begin{bmatrix} x \\ y \\ \theta \end{bmatrix}$ drives the error e towards zero.

kinematic model

We assume that the goal is at the origin of the inertial frame. In the following the position vector $[x, y, \theta]^T$ is always represented in the inertial frame. The kinematics of a differential-drive mobile robot described in the inertial frame $\{X_I, Y_I, \theta\}$ is given by:

$${}^I \begin{bmatrix} \dot{x} \\ \dot{y} \\ \dot{\theta} \end{bmatrix} = \begin{bmatrix} \cos\theta & 0 \\ \sin\theta & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} v \\ \omega \end{bmatrix}$$

where \dot{x}, \dot{y} are the linear velocities in the direction of the X_I and Y_I of the inertial frame. Let α denote the angle between x_R axis of the robot's reference frame and the vector \hat{x} connecting the center of the axle of the wheels with the final position. If $\alpha \in I_1$, where $I_1 = [-\pi/2, \pi/2]$, then consider the coordinate transformation into polar coordinates with its origin at the goal position:

$$\rho = \sqrt{\Delta x^2 + \Delta y^2}, \alpha = -\theta + \text{atan2}(\Delta y, \Delta x), \beta = -\theta - \alpha$$

This yields a system description, in the new polar coordinates, using a matrix equation

$$\begin{bmatrix} \dot{\rho} \\ \dot{\alpha} \\ \dot{\beta} \end{bmatrix} = \begin{bmatrix} -\cos\alpha & 0 \\ \frac{\sin\alpha}{\rho} & -1 \\ \frac{-\sin\alpha}{\rho} & 0 \end{bmatrix} \begin{bmatrix} v \\ \omega \end{bmatrix}$$

where ρ is the distance between the center of the robot's wheel axle and the goal position, θ denotes the angle between the X_R axis of the robot reference frame, and the X_I axis associated with the final position v and ω are the tangent and angular velocities respectively. On the other hand, if $\alpha \in I_2$, where $I_2 = (-\pi, -\pi/2] \cup (\pi/2, \pi]$ redefining the forward direction of the robot by setting $v = -v$, we obtain a system described by a matrix equation of the form

$$\begin{bmatrix} \dot{\rho} \\ \dot{\alpha} \\ \dot{\beta} \end{bmatrix} = \begin{bmatrix} \cos\alpha & 0 \\ \frac{-\sin\alpha}{\rho} & 1 \\ \frac{\sin\alpha}{\rho} & 0 \end{bmatrix} \begin{bmatrix} v \\ \omega \end{bmatrix}$$

The control law

The control signals v and ω must now be designed to drive the robot from its actual configuration to the goal position. It is obvious that there is a discontinuity for $\rho = 0$; thus the theorem of Brockett does not obstruct smooth stabilizability. If we consider now the linear control law:

$$v = k_\rho \rho, \omega = k_\alpha \alpha + k_\beta \beta$$

we get a closed-loop system described by

$$\begin{bmatrix} \dot{\rho} \\ \dot{\alpha} \\ \dot{\beta} \end{bmatrix} = \begin{bmatrix} -k_\rho \rho \cos \alpha \\ k_\rho \sin \alpha - k_\alpha \alpha - k_\beta \beta \\ -k_\rho \sin \alpha \end{bmatrix}$$

This system does not have a singularity in $\rho = 0$ and has a unique equilibrium point at $(\rho, \alpha, \beta) = (0, 0, 0)$. Thus it will drive the robot to this point, which is the goal position.

Local stability issue

It can further be shown, that the closed-loop control system is locally exponentially stable if $k_\rho > 0; k_\beta < 0; k_\alpha - k_\rho > 0$. (I'm skipping the proof cause it is tedious, but it is a page 103(physically 88) of the book)

Architectures

Notes from classes

AI has converged on a canonical operational architecture after years of implementation: historical development is of interest because it helps identify legacy code and to see barriers that led to the convergence. Operational architectures have advantages such as: The programmer is forced to apply the software engineering principles of abstraction and modularity; checklist of completeness; design specification and it can also reveal vulnerabilities.

Overall style of design or organization is called software architecture, which describes a set of architectural components and how they interact. It also provides a principled way of organizing a control system. However, in addition to providing structure, it imposes constraints on the way the control problem can be solved[Mataric].

We have different types of architectures:

- Operational architecture: What the system does at a high level and NOT HOW it does it
- Systems architecture: Describes how a system works in terms of major subsystems
- Technical architecture: It describes how a system works in terms of implementation details (even prescribing the programming language).

The last two blocks focus on implementation and software. We worry so much about organization because the 7 distinct areas of Artificial Intelligence have to be knitted together somehow. Software engineering is necessary for a successful software enterprise. Thinking about architectures is good software engineering due to the following aspects:

- **Abstraction:** Ignores details to permit focus for thinking about general organization of intelligence and is a semi-formal description
- **Modularity:** It gives us high cohesion, low coupling and supports unit testing and debugging
- **Anticipation of change, Incrementally:** How to adapt, support evolution
- **Generality:** Not re-invent the wheel each time.

The canonical robot's architecture is the following: Planner, sequencer and Reactive (or behavioral) Layer. The upper brain or cortex is used to reason over symbols about goals; the middle brain is converting sensor data into symbols; The spinal cord and lower brain are used for skills and responses. Unfortunately, in animals there is an emotional intelligence too. We can use the following to cope with that:

- user interfaces: Display, transparency of what robot is doing and thinking; natural language and gestures.

- Working in teams: Explicit multi-agent coordination often relies on social rules
- Persona we present to others: Security as what can you at your security grade can see about me; Human-robot interaction as in affective responses and natural language.

The extended canonical architecture is composed of the following 4 elements: Interaction (Persona and teaming, Deliberative loop (Reasoning over symbols), Key function (converting sensor data into symbols, Reactive loop (skill and responses). The hard part is the following: from reactive to deliberate we have two types of perception (Direct and Recognition, which impacts computer vision), and we have two different time horizons: from present to present, past and future; impact sensing, storage, as well as algorithm reasoning, projecting. It also needs a central structure to hold the symbols, history, knowledge but is tractable. From Reactive/Deliberate to interaction we need additional knowledge about "theory of mind"-beliefs, desires, intentions of the other agent and common ground.

Making things more tangible

The primitives for robot intelligence are the following: Sense, Plan, Act and Learn. The problem is that not everything intelligent requires deciding. Behavioral robotics was born with sets of SENSE-ACT couplings called behaviors that get turned on/off based on situations. There is no Plan. Hybrid architecture have can Plan and then Sense-Act, or Plan and then instantiate appropriate Sense-Act behaviors, until next step in plan. Plan requires a world model plus the actual planning algorithm. Control theory is lower level but doesn't necessarily capture it all. Reactive is tightly coupled with sensing, so very fast; many concurrent stimulus-response behaviors, strung together with simple scripting with FSA. the action is generated by sensed or internal stimulus. There is no awareness, no mission monitoring. Models are of the vehicle, not the larger world. For each layer we have different programming languages: for the interaction Layer we need Procedural, Functional, Ontological languages such as OWL; for the Deliberative Layer we need Functional languages such as LISP; while for the behavioral layer we need Procedural languages such as C, C++ or Java.

Primitives for robot intelligence

We have three primitives that we can identify for our robot: SENSE, PLAN, ACT. The book introduced another primitive, which is the LEARN. Intelligent systems in nature, like animals or simple form of lives can learn from the environment and create some memory of what happened before in order to increase the survival of themselves. How do you implement or arrange these primitives? We have three different paradigms for organising intelligence: Hierarchical, Reactive or Deliberative/reactive hybrid. These days we might have also deep learning robotics. We have two ways to describe them through the relationships between the three primitives of robotics (SENSE,PLAN,ACT) and how sensory data is managed and distributed in the system.

Hierarchical robotics

The first solution was the classical AI approach aka the hierarchical approach: SENSE - PLAN - ACT and then start over again. We have a sequence of 3 primitives. The classical planning involves predicate Logic, logical actions and planning algorithms. We also use task policies and we integrate task and motion planning. What should we put into each of this modules? for SENSE we need sensors and feature extractions; for PLAN we need to combine features into model, plan

tasking and task execution; for ACT we need motor control and actuators. The hierarchical model has a centralized sensing. The world model is a fused global data structure. It combines a priori representation, sensed info and cognitive understanding. The SENSE takes in input sensor data and gives in output sensed information. The PLAN takes in input information (sensed and/or cognitive) and gives in output directives. The ACT takes in input sensed information or directives and gives in output the actuator commands. We talked about Shakey (Hope it is not in the exam). While hierarchies have advantages, relying on a world model creates problems. We have problems on processing, particularly for control. The alternative is to create layers or hierarchies within the world model to match other subsystems. Another problem is that the world model requires extensive representation which leads to two major problems: it operates under the closed world assumption and frame problem. In practice, the implementations are planning-centric. The programmer must implement: a representation of the model of the world, a differentiable with operators, preconditions and postconditions, and different evaluation function. Strips (used for Shakey) suffered from frame problem and assumes a closed world. But remember, not everything intelligent requires "deciding" (reflexes or muscle memory such as riding a bike or driving a car).

Reactive robotics

it raises from two reasons: dissatisfaction with the results of the hierarchical paradigm and the influence of ideas from Wiener's ethology and cybernetics. We have direct relation between SENSE and ACT. We only consider the reactive intelligence of humans. Users loved it because it worked. AI people loved it but wanted to put PLAN back in. Control people hated it because couldn't rigorously prove it worked. The reactive approach didn't give a fixed models or a mathematical model to apply for different tasks. Cybernetics is a combination of control theory, information theory and biology. It tries to explain the principles of control in both animals and machines. Uses the math of feedback control systems to express natural behaviors. The emphasis is on the strong coupling between the organism and its environment. Wiener is the initiator of cybernetics. Now we are talking about turtle. The main behaviors of turtle are Thrifty, Exploration, Attraction, Aversion and Recharging. Complex behaviors arise from the combination of these simple actions. **Emerging intelligence:** the capability to exhibit a new behavior that was not coded but it arose as activation or combination of singular, simpler behaviors.

We can classify three broad categories of behaviors: Reflexive behaviors (stimulus-response. Usually hard-wired for fast response), Reactive behaviors (learned. Usually compiled down to be executed without conscious thought), Conscious behavior (requires deliberative thought).

Notes from Murphy

The term architecture is frequently used in robotics to refer to the overall style of design or organization. Architectures distill prior experience into templates for creating new intelligent robots. We have 3 types of architectures:

- *Operational architecture:* describes what the system does, or its functionality, at a high level, but not how it does it. An operational architecture specifies how the seven distinct areas of artificial intelligence, each with their own algorithms and data structures, are knitted together and can be used to determine if a design provides the intended functionality.
- *Systems architecture:* describes how a system is decomposed into major subsystems. The system architecture states the specific software systems that supply the desired intelligent

functionality and describes how they are connected. It can be used to determine if a design meets good software engineering principles, especially modularity and extensibility.

- *Technical architecture*: describes how the implementation details of the system. It specifies the algorithms and even the languages that the modules are programmed in and it can be used to determine if a design is using the most appropriate algorithms. They are in constant evolution, reflecting advances in AI.

We want to use the four general principles of software engineering:

- **Abstraction**, where the operational and systems architecture portals allow designers to ignore details in order to focus on thinking about general organization of intelligence.
- **Modularity**, where systems and technical architectures encourage the designer to think in terms of object-oriented programming. Modules should have high cohesion, where each module or object does one thing well and unrelated functions are put elsewhere, and low coupling, where modules or objects are independent. High cohesion and low coupling support unit testing and debugging.
- **Anticipation of change with incrementality**, as the systems and technical architectures provide the designer with insights whether the code is engineered to support upgrading and algorithm with a newer one and adding new capabilities without requiring the designer to rewrite the code for existing modules and possibility to introduce bugs.
- **Generality**. The operational, systems and technical architectures provide the designer with frameworks to determine if the basic organization, subsystems and implementation will allow the code to be used for other applications. In robotics, this principle is often called portability.

In addition to the four general principles of software engineering, Arkin adds two more software engineering principles for AI robotics: niche targetability and robustness. The first one captures how well the robot works for the independent application, while robustness identifies where the system is vulnerable, and how the system design intrinsically reduces that vulnerability.

Canonical AI Robotics Operational Architecture

Organizationally, the operational architecture consists of three layers, reactive, deliberative and interactive. Biological intelligence can be thought of as consisting of four major functions, reaction, deliberation, conversion of signals into symbols and interaction of the robot with other external agents. This abstraction of biological intelligence leads to three distinct purpose and styles of computing. The 3 layers not only encapsulate philosophically different categories of functionality, they have 5 different attributes which influence computing: primitives, perceptual ability, planning horizon, time scale and use of models.

Primitives: Similar to function in the LOA and ACL architectures, AI robotics views autonomous capabilities as consisting of four primitives: SENSE, PLAN, ACT, and, LEARN. If a function is taking information from the robot's sensors and producing an output useful by other functions, then that function is into the SENSE category. If a function is taking in information and producing one or more tasks for the robot to perform, that function is in the PLAN category. Functions which produce output commands to motor actuators fall into ACT. Intelligence is generally viewe as a process of SENSE, PLAN and ACT. LEARN is an important mechanism by which an agent maximizes its chances for success, but it transcends the other 3 primitives; an agent can learn to

sense or plan better, acquire more acts or skills, or even learn for a specific mission what to sense, what to plan for and how to act.

Perpetual ability: AI robotics view the perception needed for a particular capability as being either direct or requiring recognition. This divides capabilities into those that can work directly with incoming stimuli and those requiring a conversion of the stimulus into a symbol.

Planning horizon: AI robotics views a capability as also having a planning horizon of either present;past,present; or future,past,present. The planning horizon constrains the choice of data structures, as well algorithms.

Time-scale: AI robotics also considers the time-scale of a function,i.e., whether it need to be very fast, fast, or can be relatively slow. Time-scale helps determine asynchronous operation of the software components.

Reactive(behavioral) layer

Reactive functionality corresponds to the reactive loop in the central nervous system. This loop is associated with functionality that occurs in the spinal cord and lower brain, especially responses and skills based on motor memory. The responses and skills are patterns of action generally referred to as behaviors. Behavioral functionality is sufficient for animal intelligence. In robotics, it consists of functionality constructed from the SENSE and ACT primitives with no PLAN. SENSE and ACT are tightly coupled into constructs called behaviors which produce capabilities. The perceptual ability of the behaviors is called direct perception,. While a behavior does not create and maintain a global world model, a behavior may have a local world model, which serves as short term memory. There is no planning, and thus planning horizon is the present. The time-scale for the functions is very fast. Stimulus-response behaviors typically have a fast update cycle on the order of 15-30 cycles per second, matching control need and sensor update rates.

Deliberative Layer

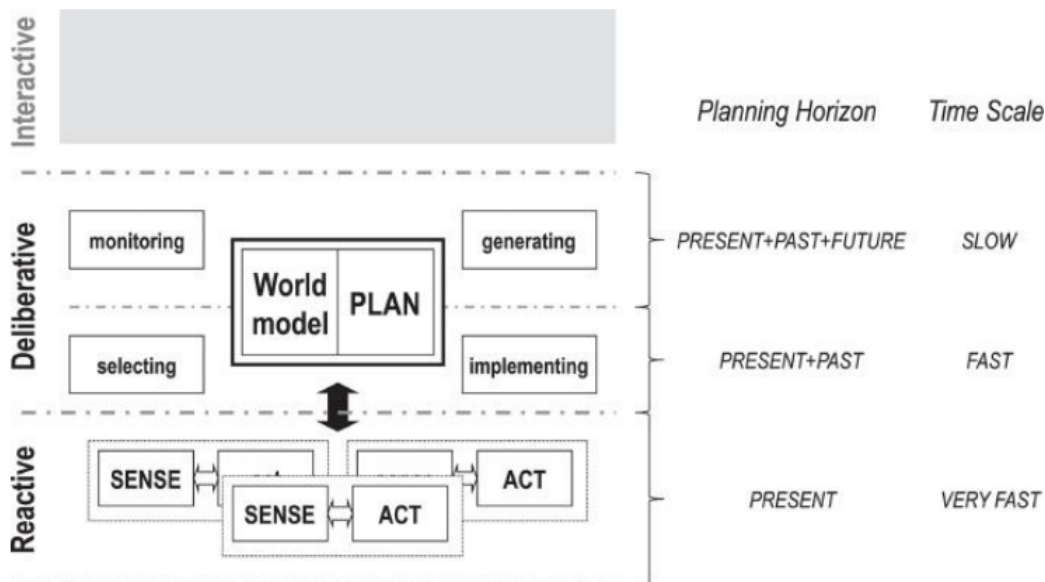
Deliberative functionality corresponds to the cognitive loop associated with the cortex in the brain. The brain independently takes as input the same signals used by other ventral nervous system and adds additional sensor processing to make conscious decisions. The cognitive loop can modify the reactive loop, either instantiate new actions that would be carried out by reactive functionality or to modify them. Deliberative functionality is sufficient for the more sophisticated problem solving and reasoning aspects of intelligence as well as planning. The two loops differ in processing speed and the content of what is being processed. Reflexes and motor skills are very fast, while deliberating about a problem may be slow. The deliberative layer hosts the PLAN activities of the robot, has a Planner that generates a plan using a World Model and the instantiates appropriate SENSE-ACT behaviors in the reactive layer to execute the plan. These behaviors run until the next step of the plan is reached and a new set of behaviors is instantiated or the deliberative monitoring detects a problem with the plan. The perceptual ability of the behaviors is recognition, where the robot builds up global or persistent world models. The planning horizon uses information from the past and the present to project consequences of a plan into the future. The time it takes in the deliberative layer may update with a frequency ranging from on the order of 15 cycles per second to update the World model to several minutes to construct complex plans. The deliberative layer is divided into 2 sub-layers connected by the world model that contain the four deliberative functions:

- *Generating plans* which corresponds to planning, reasoning, and problem solving in AI

- *Selecting* specific resources to accomplish the plan, which corresponds to planning, resource allocation, and knowledge representation of capabilities in AI
- *Implementing* the plan, which corresponds to execution
- *Monitoring* the execution of the plan to determine if it is meeting the goal, learning what is normal, and anticipating potential failures, which corresponds to planning and reasoning in AI.

Interactive layer

Interactive functionality is need for interaction with other agents, either people, other robots, or software agents. Interaction may be regulated by reactive or cognitive loops. Displays of interactive functionality essentially surround the person and produce their persona. The interactive layer is essentially a wrapper on individual robot programming that enables the robot to work with other agents. The interactive layer may be programmed in procedural, functional, or ontological languages, such as OWL.



Other Operational Architectures

Levels of Automation (LOA)

- *Levels of Autonomy*: while biological intelligence provides insight into the general organization of intelligence, the human-machine systems community offers a different viewpoint. The human-machine system community has looked at automation for space exploration, control of processing plants, and autopilots in aircraft. The community has traditionally labeled the state of automation at any given time in terms of what functions a human has currently delegated to the computer. This classification of the state in a system naturally leads to hierarchies of where the machine has responsibility for more of the functions, and thus would be considered more autonomous. These taxonomic hierarchies are often used as operational architectures or as de facto measure of how automated a system is, where level B is more

automated than level A. The hierarchies are generally referred to as levels of automation or levels of autonomy, as automation and autonomy are often used as synonyms.

- *Traded, or Shared, Control:* A function can be performed by a human, a computer, both working together, or portion can be performed by either a human or computer and then the remaining portions by the other agent. In the first case we talk about shared control, while in the latter we talk about traded control. The state of automation for a process is defined by whether the human or the machine is performing the function. No particular levels of automation is accepted by all researchers, with different experts arguing over the functions and the ordering of the hierarchy. However, the different variants of levels of automation all share the same tenet that there are a set of functions and the more functions delegated to a computer means higher automation. The resulting levels of automation approach is often used as an operational architecture, with the expectation that an intelligent robot would be built incrementally in the layers. An unintended consequence of using levels of automation as an operational architecture is that it implies full automation as the design goal, rather than matching the appropriate level of automation for the mission and robot's capabilities. The major advantage of using the levels of automation organization is that the four functions and levels offer more specific modularity than the three layers in biological intelligence. The major disadvantage if this organization is the levels of autonomy were initially intended to be a set of definitions for labeling the current state of an active capability, not an operational architecture for coordinating all capabilities. There are other two disadvantages: treating the levels of autonomy taxonomy as an operational architecture results in a system that considers only the four deliberative functions, ignoring reactive behaviors; and levels of autonomy taxonomy considers the mode of interaction with the human supervisor but does not consider interaction with other agents or nonsupervisory interactions.

Levels of Initiative

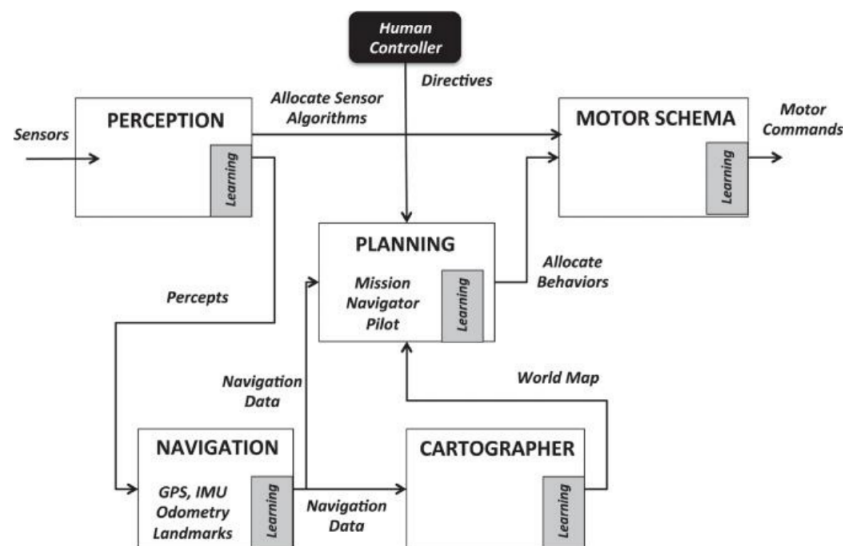
Another style of operational architecture is based on the amount of initiative that a robot is given to conduct a task. In this case, autonomy is based on the political connotation of autonomy as self-governing rather than the mechanical connotation we used previously. Unlike the other operational architectures, initiative is conceptualized by roles rather than by functions. The five levels of initiative are:

- *No autonomy:* The robot follows rigid programming to perform a task or achieve a goal. This level is similar to stimulus-response behaviors in a biological architecture and the 10 levels in the levels of automation architecture.
- *Process autonomy:* The robot can choose the algorithm or process to achieve its task goals. This choice is similar to the selecting function in a levels of automation architecture.
- *System-state autonomy:* The robot can generate and choose between options to achieve its goal. This level is similar to the generating function in a levels of automation architecture but implies increased unpredictability and nondeterminism.
- *Intentional autonomy:* The robot can change its goals. This functionality introduces the notion of being aware of others and explicitly interacting with them, similar to the interface layer in biologically-inspired architecture.

- *Constraint autonomy*: The robot can create its own roles and goals, which may involve relaxing constraints on its goals or how it accomplishes its goals. It is still subject to bounded rationality.

Five Subsystems in Systems Architectures

Operational architectures focus on the general style, while systems architectures focus on the general components. The software for an intelligent robot normally has at least five subsystems, which are encapsulated as object libraries or similar reusable programming repositories. System architectural design encourages designers to think in terms of creating libraries of algorithms and data structures for each subsystem. These libraries serve as general clearinghouses from which the designer can pick specific functions for a particular robot in order to customize it for a new application. The subsystems are the followings:



- **Planning:** is associated with generating and monitoring overall mission objectives and passing along the geospatial component of those objectives to the navigation subsystem, selecting the motor and perceptual resources, implementing or instantiating those resources, and monitoring the performance of the mission. This subsystem is host to the numerous classical planning and resource allocation algorithms.
- **Cartographer:** also known as the world model or world map, is associated with the construction and maintenance of knowledge representations about the world. For navigation-oriented robots, this subsystem is the key data structure that enables generating, monitoring, selecting, and implementing actions. These knowledge representations are often geospatial maps but can include more abstract symbolic information, such as beliefs about the state of the world or the intent of other agents. In more deliberative robots, this subsystem bridges the gap between signals and symbols. In interactive robots, the subsystem also monitors the robot's mental models of itself and the beliefs, desires, and intents of other agents.
- **Motor Schema:** also known as Motor Schema Library, or more generally as the Behaviors subsystem, is associated with selecting the best motor routines and implementing actions. This subsystem contains the functions that connect the deliberative and reactive functions

with the requisite guidance, navigation, and control algorithms for the actuators and effectors. This subsystem typically does not have deliberative functionality but is rather about execution. The algorithms in the motor schema subsystem may be programmed to monitor their execution within a control loop.

- **Perception:** also known as the Sensing Perceptual Schema, or Perceptual Schema Library, is associated with selecting the best sensors for the motor actions and implementing sensor processing algorithms. This subsystem contains the functions that control the sensors and extract the percept, that is, what is perceived. This includes algorithms that extract a special type of percept called affordances or those that perform object recognition and image understanding.

The subsystems are not independent and do not represent a sequence of programming actions. Instead, they are essentially the major object or classes in object-oriented programming that reflect software engineering decisions about how to group similar functions and data. The list of subsystems is not necessarily complete. The five common subsystems highlight the historical focus within the robotics community on autonomous navigation of a platform and mapping. The five subsystems above have no clear component for interacting with other agents or for manipulation. An intelligent robot in healthcare, entertainment, or even driving in city traffic would likely have additional subsystems.

Three System Architecture Paradigms

Historically, system architectures for AI robotics fall into three categories called paradigms: hierarchical, reactive, or hybrid deliberative/reactive. These paradigms arrange the flow of data and control within the robot. This data and control flow for each paradigm can be uniquely described using two traits: the pattern of interaction between the three robot primitives of SENSE, PLAN and ACT, and the route of sensing.

Interaction between primitives

Intelligent robots use one of three cyclic patterns of interaction between the primitives. One is to SENSE, then PLAN, and finally ACT, at which point acting changes the state of the world which lead to a new iteration of SENSE, PLAN and ACT. We also gave reactive animal intelligence, which uses a pattern of SENSE and ACT coupled into behaviors with no PLAN. Neurophysiological models favor a different pattern of intelligence where higher brain functions PLAN asynchronously while SENSE and ACT behaviors both carry out plans and react to the environment.

Sensing Route

We have three routes:

- **Local sensing:** it goes directly from the sensor to the behavior or function using the sensor data. The receiving function is responsible for any transformations or conversions of the raw data into a suitable data structure. The transformations or conversions stay local (within the programming scope of that function). Local sensing can be a one-to-one or a one-to-many mapping, where the data from one sensor can go to one or more independent functions.

- **Global sensing:** is how data originating from all the sensors is transmitted to a function that transforms and fuses the data into a global world model or unified data structure. Global sensing can be thought of as a many-to-one mapping, where the data from many sensors go to one function.
- **Hybrid sensing:** is a combination, where the same data from a sensor may go to one or more functions that perform local transformations and to a global sensing function. Hybrid sensing can be thought of as a many-to-many mapping, where the data from many sensors can go to many independent functions.

Hierarchical System Architecture Paradigm

Hierarchical systems architectures organize the five subsystems to support a SENSE-PLAN-ACT pattern that relies on a global sensing route. Hierarchies are a natural way to organize functionality and they can be computationally efficient because they can reduce computation by specifying the frame and defining a closed world. They are used for implementations where the mission or application is well-understood and further additions of capabilities or major upgrades are not expected. They are often used where there is little expectation of reuse of code or later expansion of functionality. Thus the goal is to optimize the programming.

The SENSE phase collects data from all sensors using routines from the perception subsystem and then fuses sensor data into a World Model using routines from the cartographer subsystem. It fuses the current readings with older readings and any a priori knowledge, such as maps. The PLAN phase applies routines from the planner subsystem to the World Model to determine any changes to the missions or constraints and the plans navigational routes and movements using routines from the navigational subsystem. The output is a set of actions. The ACT phase executes the actions using routines from the motor schema subsystem. Action may also involve the perception subsystem if sensors have actuators to be moved. The SENSE-PLAN-ACT sequence executes in a continuous loop.

Reactive Systems Paradigm

The Reactive Systems use only two subsystems, Perception and Motor Schema, of the five canonical subsystems to support SENSE-ACT pattern that relies on a local sensing route. Programming by behavior has a number of advantages, most of them consistent with good software engineering principles. Behaviors are inherently modular and easy to test in isolation from the system. Behaviors also support incremental expansion of the capabilities of a robot. A robot becomes more intelligent by having more behaviors. The behavioral decomposition results in an implementation that works in real-time and is usually computationally inexpensive. Generally, the reaction speeds of a reactive robot are equivalent to stimulus-response times in animal, although if the technical architecture choices are implemented poorly, then a reactive implementation can be slow. Purely reactive robots are useful for simple autonomous systems, such as Roomba, robot mowers and so on. They can be used as well to add autonomous capabilities to an existing human-machine system where a rapid response is important.

The flow of a typical reactive architecture uses multiple instances of the Perception and Action subsystems in parallel. The SENSE phase collects data from all sensors using routines from the Perception subsystem and directly supplies them to the Motor Schemas which produce Actions. There may be multiple SENSE-ACT pairs and they run concurrently and independently. Each behavior (or SENSE-ACT pair) is a local loop and may have a different computation time, and thus the behaviors work asynchronously. Sensing is routed to each behavior without any filtering

or processing and each behavior extracts the direct perception that is unique to it. A sensor can supply the same sensor reading simultaneously to different behaviors, and then each behavior compute a different percept. The sensor processing is local to each behavior with no global model of the world. There is no PLAN phase as behaviors produce actions only if the sensing produces a reaction.

A robot constructed only with the reactive layer is often called either a reactive or behavioral robot. A behavioral robot has at least 3 advantages: direct perception is usually simple to program; the behaviors are highly modular; It can work well in cooperation with a human in telesystems, because the robot can react to stimuli to trigger panic behaviors or guarded motion behaviors. A reactive robot has 3 disadvantages: the inability to PLAN leads to the "fly-at-the-window" effect; the emergence of the overall behavior from individual behaviors using direct perception makes it difficult to precisely predict what the robot will do. The behavior may not be optimal because it is based on local, instantaneous sensing rather than on a plan generated from a global model.

Hybrid Deliberative/Reactive Systems Paradigm

This kind of architecture organizes the five subsystems to support a "PLAN, then SENSE-ACT" pattern coupled with hybrid sensing. The idea of this pattern evolved both from biomimicry and from two software engineering assumptions: planning covers a long time horizon and requires global knowledge; planning and sensor fusion algorithms are computationally expensive. The world model is constructed by the process independent of the behavior-specific sensing. Hybrid systems use hybrid sensing routes. The cartographer or model making processes can access not only the same sensors as the behaviors but also the results of local sensing. The cartographer can also have sensors that are dedicated to providing observations which are useful for world modeling but are not used for any active behavior. Hybrid systems are the most flexible and have become the default systems architecture for researchers and for applications where the robot's functionality is expected to be modified or extended in the future.

Execution Approval and Task Execution

The many facets of deliberation and reaction pose the concern as to whether we have algorithms and coordination methods to actually build a fully autonomous robot that can work in the open world. One strategy is to build the robot but have a human supervisor check the robot's planned actions before it executes them; this strategy is known as execution approval and task rehearsal. In the case of execution approval, the cognitive agent (human or computer) considers the next step and decides whether it is safe. Task rehearsal verifies the plan by considering an entire sequence of actions.

Behaviors

Even though it seems reasonable to explore biological and cognitive sciences for insights into intelligence, how can we compare such different systems: carbon and silicon "life" forms? One powerful means of conceptualizing the different systems is to think of an abstract intelligent system as an agent. In object-oriented programming terms, the agent is the superclass, and the classes of persona, animal and robot agents are derived from it. one helpful way of seeing correspondences is to decide on the level at which these entities have something in common. The set of levels of commonalities leads to a **computational theory**. The levels in a computational theory can be greatly simplified as:

- **Level 1: Existence proof of what can/should be done.** At Level 1, agents can share commonality of purpose or functionality.
- **Level 2: Decomposition of "what" into inputs, outputs, and transformations.** This level can be thought of as creating a flow chart of black boxes to actually accomplish the "what" from Level 1. Each box represents a transformation from input to output. At Level 2, agents can exhibit common processes.
- **Level 3: How to implement the process.** This level of computational theory focuses on describing how each transformation is implemented, like a neural network or an algorithm. At Level 3, agents may have little or no commonality in the actual physical implementation of the functionality specified in Level 1.

More precisely, Level 1 corresponds to a phenomenon that people can do. The second level refines the description of the phenomenon into a series of transformations that would describe what sensor data enters the brain and what region of the brain are activated and produce outputs leading to the ultimate output. The third level describes the activity of each region of the brain in terms of the specific types of neurons and how they are connected. it should be clear that Levels 1 and 2 are abstract enough to apply to any agent. It is only Level 3 that the differences between a robotic agent and a biological agent really emerge.

Rana Computatrix

Starting with Level 1 of the computational theory approach, toads and frogs have visually directed behaviors: feed and flee. Level 2 requires expressing the feed and flee behaviors as inputs, outputs, and transformations. From biology studies, frogs can be characterized as being able to detect only two visual inputs: small moving blobs on their visual field or large moving blobs. The small moving blobs are inputs to the feeding behavior, which produce the output. The output is that the frog orients itself towards the blob and then snaps at it. large moving blobs are inputs to the fleeing behavior, causing the frog to hop away.

At Level 2, the stimulus output is called a **percept**.

Exploring the Level 3, they implemented the taxis behavior as a vector field: rana computatrix would literally feel an attractive force along the direction of the fly. The relative direction and intensity of the stimulus was represented as a vector. The direction indicated where the rana had to turn and the magnitude indicated the strength that should be applied to snapping. At Level 3 the percept and the output are more precisely defined, though still not expressed in a true software engineering format. In the case of the rana computatrix implementation, if it gets more than an input, it would sum the vectors, resulting in the generation of a new vector in-between of the original inputs.

Animal Behavior

Rana computatrix illustrates how a computational theory spans biological and artificial intelligence. A **Behavior** is a mapping of sensory inputs to a pattern of motor actions which then are used to achieve a task. We have different kind of behaviors: **REFLEXIVE BEHAVIOR, STIMULUS-RESPONSE, REACTIVE BEHAVIOR, CONSCIOUS BEHAVIOR**. behaviors can then be divided into 3 broad categories: **Reactive behaviors, Reflexive behaviors and conscious behaviors**. Reflexive behaviors are stimulus-response. Reactive behaviors are learned and then consolidated to where they can be executed without conscious thought. Conscious behaviors are deliberative.

Reflexive Behaviors

They imply no need for cognition: if you sense it, you do it. For a robot it means a hardwired response, eliminating computation and guaranteed to be fast. Reflexive behaviors can then be divided into 3 categories:

- **Reflexes**: where the response lasts only as long as the stimulus, and the response is proportional to the intensity of the stimulus.
- **Taxes**: where the response is to move to a particular orientation
- **Fixed-action patterns**: where the response continues for a longer duration than the stimulus.

These categories are not mutually exclusive.

Behavioral Coordination

The overall response of robot ρ is a function of the behaviors B , the gain of each behavior G , and the **coordination function** C . The coordination function C takes the output of multiple behaviors and produces one response or output. More formally $\rho = C(G \times B(S))$. There are 3 categories of coordination functions. A set of behaviors can be either **concurrent**, that is, all are acting at the same time, or they can be acting as a **sequence**. If the behaviors are concurrent, there are two broad categories of algorithms for coordinating the set of behaviors. One set is the **cooperating methods**, where the coordination function blends the output of the individual behaviors into a single output. The most common methods are vector summation. The other category of concurrent algorithms is **competing methods**, where the individual behaviors undergo arbitration. The most common methods of competition are subsumption and voting. A group of behaviors that form a sequence or macro can be grounded into an **abstract behavior** that may use a finite state automata to explicitly coordinate behaviors while any other behaviors, such as avoiding obstacles, might use potential fields.

Potential Fields

The most common method for cooperating behaviors is to represent the output as a vector. Often the vector is produced by the motor schema using a potential field representation of possible actions, typically referred to as a **potential field methodology**. They always use vectors to represent the action output of behaviors and vector summation to combine vectors from different behaviors to produce an emergent behavior.

Visualizing potential fields

We need to represent the motor action of a behavior as a potential field, which is an array, or field, of vectors. It has a magnitude (real number from 0.0 to 1, if normalized, otherwise any real number) and a direction. The array represents a region of the space, usually 2 dimensional. We can divide the map into squares, where each array represents a square of the space. Perceivable objects in the world would exert a force field on the surrounding space. The vector in each element represents the force, both the direction to turn and the magnitude or velocity to head in that direction. Potential fields are continuous because it does not matter how small the element is. The field represents what the robot should do based on whether the robot perceives an obstacle. The field is not concerned with how the robot came to be so close to the obstacle; the robot feels the same force if it happens to move within range or if it was just sitting there and someone put a hand next to the robot. One way of thinking about potential fields is to imagine a force field acting on the robot.

There are five basic potential fields, or primitives, which can be combined to build more complex fields:

- **Uniform field:** the robot would feel the same force no matter where it was. No matter where it got set down and at what orientation, it would feel a need to align itself with the direction of the arrow and to move in that direction at a velocity proportional to the length of the arrow. (go to the direction n°).
- **Perpendicular field:** where the robot is oriented perpendicular to some object or wall or border. It can be pointed away or toward an object.
- **Attractive field:** the center of the field represents an object that is exerting an attraction to the robot. Wherever the robot is, the robot feels a force relative to the object. Attractive fields are to represent cases where the agent is attracted to light, food or goal.
- **Repulsive field:** opposite of an attractive field. Commonly associated with obstacles or things the agent should avoid. The closer the robot is to the object, the stronger the repulsive force pushes the robot 180 degrees away from it.
- **Tangential field:** the field is tangent to the object. Tangential field can spin either clockwise or counterclockwise. These fields are useful for directing a robot to go around an obstacle or having a robot investigate something.

Magnitude profiles

The **magnitude profile** is the way the magnitude of the vectors change. Mathematically the field can be presented with polar coordinates, and the center of the field being the origin (0,0).. Magnitude profiles resolve the problem of constant magnitude. They also make it possible for a robot designer to represent flexibility and to create interesting responses. Imagine the following scenario: if the robot is far away from the object, it will turn and move quickly towards it and then slow up to keep from overshooting and hitting the object. Mathematically, this is called a **linear drop off**, since the rate at which the magnitude of the vectors drops off can be plotted as a straight line. This profile also shares the problem of the constant magnitude profile in the sharp transition to 0.0 velocity. Therefore, another profile might be used to capture the need for a stronger reaction but with more of a taper. Such profile is an **exponential drop off** function, where the drop off magnitude is proportional to the square of the distance.

Potential fields and Perception

The force of the potential field at any given point was a function of both the relative distance between the robot and an object and the magnitude profile. The strength of a potential field can be a function of the stimulus, regardless of distance.

Programming a Single Potential Field

They are actually easy to program, especially since the fields are egocentric to the robot. The robot computes the effect of the potential field, usually as a straight line, at every update, with no memory of where it was previously or where it has moved from.

A primitive potential field is usually represented by a single function. The vector impacting the robot is computed at each update. Considering the case of a robot with a single range sensor facing forward. The designer has decided that a repulsive field with a linear drop off is appropriate.

The formula is:

$$V_{direction} = -180, \quad V_{magnitude} = \begin{cases} \frac{(D-d)}{D} & \text{for } d \leq D \\ 0 & \text{otherwise} \end{cases}$$

where D is the maximum range of the field's effect, or the maximum distance at which the robot can detect the obstacle, and d is the distance of the robot to the obstacle.

Combination of Fields and Behaviors

The first attribute of a true potential field methodology is that it requires all behaviors to be implemented as potential fields. The second attribute is that it combines behaviors by **vector summation**. A robot will generally have forces from multiple behaviors acting on it concurrently. One problem is that the sum may be 0.0, leading to the **local minima problem**, i.e. the potential field has a minima, or valley, that traps the robot. (If you want more details look at the Murphy book).

Davantages and Disadvantages

The potential field is a continuous representation that is easy to visualize over a large region of space. As a result, it is easier for the designer to visualize the robot's overall behavior. The potential fields can be parametrized: their range of influence can be limited, and any continuous function can express the change in magnitude over distance. Furthermore, a two-dimensional field can usually be extended to a 3-dimensional field, and so behaviors developed for 2D will work for 3D.

Building a reactive system with potential fields is now without disadvantages. The most commonly cited problem with fields is that multiple fields can sum to a vector with 0 magnitude. There are many elegant solutions to this problem: one of the earliest was to always have a motor schema producing vectors with a small magnitude from **random noise**. Another solution is the **navigation templates**. The avoid behavior receives, as input, the vector summed from the other behaviors. This strategic behavior represents the direction the robot would take if there were no obstacles nearby. If the robot has a strategic vector, that vector gives a clue as to whether an obstacle would be passed on the right or left.

A third solution to the local minima problem is to express the fields as **harmonic functions**. They are guaranteed not to have local minima of 0. This is more expensive as the entire field has to be computed for large areas, not just for the vectors acting on the robot.

Competing Methods: Subsumption

You at 184.